

# CS315: DATABASE SYSTEMS INDEXING

**Arnab Bhattacharya**

`arnabb@cse.iitk.ac.in`

Computer Science and Engineering,  
Indian Institute of Technology, Kanpur

<http://web.cse.iitk.ac.in/~cs315/>

2<sup>nd</sup> semester, 2022-23

Tue 10:30-11:45, Thu 12:00-13:15

- Indexing is used to speed up search
- A **search key** is used
- An **index file** consists of records or **index entries** which has two fields
  - 1 Search key: Attribute that is used for searching
  - 2 Pointer to the entire object or tuple
- Index files should be smaller than data files

- Indexing is used to speed up search
- A **search key** is used
- An **index file** consists of records or **index entries** which has two fields
  - 1 Search key: Attribute that is used for searching
  - 2 Pointer to the entire object or tuple
- Index files should be smaller than data files
- Index evaluation metrics
  - Search time
  - Modification overhead
  - Space overhead

- Indexing is used to speed up search
- A **search key** is used
- An **index file** consists of records or **index entries** which has two fields
  - 1 Search key: Attribute that is used for searching
  - 2 Pointer to the entire object or tuple
- Index files should be smaller than data files
- Index evaluation metrics
  - Search time
  - Modification overhead
  - Space overhead
- Two basic types of indices
  - 1 **Ordered index**: search keys are organized according to some order
  - 2 **Hash index**: search keys are organized according to a hash function

# Static Hashing

- A **hash function** maps a key to a bucket
- A **bucket** is a unit of storage
- It is typically a disk block
- A key may need to be searched sequentially inside a bucket
- Results in **hash file organization**
- Example:  $\text{mod } n$  where  $n$  is the number of buckets

# Hash Function

- Two important qualities of an ideal hash function
- **Uniform**: Total number of keys from the domain is spread uniformly over all the buckets
- **Random**: Number of keys in each bucket is same irrespective of the actual distribution of keys

# Bucket Overflow

- Buckets may overflow at runtime due to
  - Skew in distribution of keys
  - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**

# Bucket Overflow

- Buckets may overflow at runtime due to
  - Skew in distribution of keys
  - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing:**



# Bucket Overflow

- Buckets may overflow at runtime due to
  - Skew in distribution of keys
  - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing**: Overflow buckets are linked together
  - Also called **separate chaining** or **chaining**

# Bucket Overflow

- Buckets may overflow at runtime due to
  - Skew in distribution of keys
  - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing**: Overflow buckets are linked together
  - Also called **separate chaining** or **chaining**
- **Open addressing**:

# Bucket Overflow

- Buckets may overflow at runtime due to
  - Skew in distribution of keys
  - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing**: Overflow buckets are linked together
  - Also called **separate chaining** or **chaining**
- **Open addressing**: Excess keys assigned to some other bucket and total number of buckets is kept fixed
  - Address at  $i^{\text{th}}$  attempt for key  $k$  is  $h(k, i) = (h(k) + o_i \cdot i) \bmod m$
  - **Linear probing**:

# Bucket Overflow

- Buckets may overflow at runtime due to
  - Skew in distribution of keys
  - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing**: Overflow buckets are linked together
  - Also called **separate chaining** or **chaining**
- **Open addressing**: Excess keys assigned to some other bucket and total number of buckets is kept fixed
  - Address at  $i^{\text{th}}$  attempt for key  $k$  is  $h(k, i) = (h(k) + o_i \cdot i) \bmod m$
  - **Linear probing**: Interval between buckets is fixed
    - $h(k, i) = (h(k) + c \cdot i) \bmod m$
  - **Quadratic probing**:

# Bucket Overflow

- Buckets may overflow at runtime due to
  - Skew in distribution of keys
  - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing**: Overflow buckets are linked together
  - Also called **separate chaining** or **chaining**
- **Open addressing**: Excess keys assigned to some other bucket and total number of buckets is kept fixed
  - Address at  $i^{\text{th}}$  attempt for key  $k$  is  $h(k, i) = (h(k) + o_i \cdot i) \bmod m$
  - **Linear probing**: Interval between buckets is fixed
    - $h(k, i) = (h(k) + c \cdot i) \bmod m$
  - **Quadratic probing**: Interval between buckets increase linearly
    - $h(k, i) = (h(k) + c \cdot i \cdot i) \bmod m$
  - **Double hashing**:

# Bucket Overflow

- Buckets may overflow at runtime due to
  - Skew in distribution of keys
  - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing**: Overflow buckets are linked together
  - Also called **separate chaining** or **chaining**
- **Open addressing**: Excess keys assigned to some other bucket and total number of buckets is kept fixed
  - Address at  $i^{\text{th}}$  attempt for key  $k$  is  $h(k, i) = (h(k) + o_i \cdot i) \bmod m$
  - **Linear probing**: Interval between buckets is fixed
    - $h(k, i) = (h(k) + c \cdot i) \bmod m$
  - **Quadratic probing**: Interval between buckets increase linearly
    - $h(k, i) = (h(k) + c \cdot i \cdot i) \bmod m$
  - **Double hashing**: Interval computed using another hash function  $h'$ 
    - $h(k, i) = (h(k) + h'(k) \cdot i) \bmod m$

# Bucket Overflow

- Buckets may overflow at runtime due to
  - Skew in distribution of keys
  - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing**: Overflow buckets are linked together
  - Also called **separate chaining** or **chaining**
- **Open addressing**: Excess keys assigned to some other bucket and total number of buckets is kept fixed
  - Address at  $i^{\text{th}}$  attempt for key  $k$  is  $h(k, i) = (h(k) + o_i \cdot i) \bmod m$
  - **Linear probing**: Interval between buckets is fixed
    - $h(k, i) = (h(k) + c \cdot i) \bmod m$
  - **Quadratic probing**: Interval between buckets increase linearly
    - $h(k, i) = (h(k) + c \cdot i \cdot i) \bmod m$
  - **Double hashing**: Interval computed using another hash function  $h'$ 
    - $h(k, i) = (h(k) + h'(k) \cdot i) \bmod m$
- Changing size of a database is a problem
- Periodic re-hashing is the only solution
- **Dynamic hashing**:  $h$  changes dynamically but deterministically

# Dynamic Hashing

- Dynamic hashing



# Dynamic Hashing

- Dynamic hashing
- Organize overflow buckets as binary trees
- $m$  binary trees for  $m$  primary pages

# Dynamic Hashing

- **Dynamic hashing**
- Organize overflow buckets as binary trees
- $m$  binary trees for  $m$  primary pages
- $h_0(k)$  produces index of primary page
- Particular access structure for binary trees

# Dynamic Hashing

- **Dynamic hashing**
- Organize overflow buckets as binary trees
- $m$  binary trees for  $m$  primary pages
- $h_0(k)$  produces index of primary page
- Particular access structure for binary trees
- Family of functions  $g(k) = \{h_1(k), \dots, h_i(k), \dots\}$
- Each  $h_i(k)$  produces a *bit*
- At level  $i$ , if  $h_i(k) = 0$ , take left branch, otherwise right branch

# Dynamic Hashing

- **Dynamic hashing**
- Organize overflow buckets as binary trees
- $m$  binary trees for  $m$  primary pages
- $h_0(k)$  produces index of primary page
- Particular access structure for binary trees
- Family of functions  $g(k) = \{h_1(k), \dots, h_i(k), \dots\}$
- Each  $h_i(k)$  produces a *bit*
- At level  $i$ , if  $h_i(k) = 0$ , take left branch, otherwise right branch
- Example: bit representation



# Ordered Index

- **Index sequential file**: ordered sequential file with an index

# Ordered Index

- **Index sequential file**: ordered sequential file with an index
- **Primary index** or **clustering index**: index whose search key specifies the sequential order of the file
  - Generally primary key
- **Secondary index** or **non-clustering index**: any other index

# Ordered Index

- **Index sequential file**: ordered sequential file with an index
- **Primary index** or **clustering index**: index whose search key specifies the sequential order of the file
  - Generally primary key
- **Secondary index** or **non-clustering index**: any other index
- **Dense index**: index record appears for every key
  - Secondary index must be dense
- **Sparse index**: index record appears for only some keys
  - Records must be sequentially ordered by key
  - To locate  $k$ , find largest key  $< k$ , and then sequential search from there

# Ordered Index

- **Index sequential file**: ordered sequential file with an index
- **Primary index** or **clustering index**: index whose search key specifies the sequential order of the file
  - Generally primary key
- **Secondary index** or **non-clustering index**: any other index
- **Dense index**: index record appears for every key
  - Secondary index must be dense
- **Sparse index**: index record appears for only some keys
  - Records must be sequentially ordered by key
  - To locate  $k$ , find largest key  $< k$ , and then sequential search from there
  - Deletion: Replaced by next key
  - Insertion: Inserted only if entries exceed a block
- Sparse index requires less space and less maintenance but more searching time



# Ordered Index

- **Index sequential file**: ordered sequential file with an index
- **Primary index** or **clustering index**: index whose search key specifies the sequential order of the file
  - Generally primary key
- **Secondary index** or **non-clustering index**: any other index
- **Dense index**: index record appears for every key
  - Secondary index must be dense
- **Sparse index**: index record appears for only some keys
  - Records must be sequentially ordered by key
  - To locate  $k$ , find largest key  $< k$ , and then sequential search from there
  - Deletion: Replaced by next key
  - Insertion: Inserted only if entries exceed a block
- Sparse index requires less space and less maintenance but more searching time
- **Multilevel index**: primary index does not fit in memory
  - **Outer index**: Sparse primary index
  - **Inner index**: Dense primary index file



- Balanced hierarchical data structure

# B-Tree

- Balanced hierarchical data structure
- Keys (and associated objects) are in secondary storage, i.e., disk

# B-Tree

- Balanced hierarchical data structure
- Keys (and associated objects) are in secondary storage, i.e., disk
- A **B-tree** of order  $\Theta$  has the following properties:
  - 1 Leaf nodes are in same level, i.e., the tree is balanced
  - 2 Root has at least 1 key
  - 3 Other internal nodes have between  $\Theta$  and  $2\Theta$  keys
  - 4 An internal node with  $k$  keys have  $k + 1$  children
  - 5 Child pointers in leaf nodes are null
- Branching factor is between  $\Theta + 1$  and  $2\Theta + 1$
- Pointer to the object corresponding to a key is stored alongside




# B+-Tree


- **B+-tree** is the most important variety of B-tree
- Internal nodes do *not* contain pointers to objects
- Data is stored *only* at the leaves

# B+-Tree

- **B+-tree** is the most important variety of B-tree
- Internal nodes do *not* contain pointers to objects
- Data is stored *only* at the leaves
- Often siblings are connected by pointers to avoid parent traversal




B-tree node




B+-tree node

# B+-Tree

- **B+-tree** is the most important variety of B-tree
- Internal nodes do *not* contain pointers to objects
- Data is stored *only* at the leaves
- Often siblings are connected by pointers to avoid parent traversal



B-tree node



B+-tree node

- More keys can fit in a B+-tree
- Height may be less

# Order

- Order mainly determined by disk page size



# Order

- Order mainly determined by disk page size
- Disk page capacity is  $C$  bytes
- Key consumes  $\gamma$  bytes
- Pointer to key (or object) requires  $\eta$  bytes

# Order

- Order mainly determined by disk page size
- Disk page capacity is  $C$  bytes
- Key consumes  $\gamma$  bytes
- Pointer to key (or object) requires  $\eta$  bytes
- For B+-tree

# Order

- Order mainly determined by disk page size
- Disk page capacity is  $C$  bytes
- Key consumes  $\gamma$  bytes
- Pointer to key (or object) requires  $\eta$  bytes
- For B+-tree

$$C \geq 2.\Theta.\gamma + (2.\Theta + 1).\eta$$
$$\Rightarrow \Theta_{\text{B+-tree}} = \left\lfloor \frac{C - \eta}{2(\gamma + \eta)} \right\rfloor$$

# Order

- Order mainly determined by disk page size
- Disk page capacity is  $C$  bytes
- Key consumes  $\gamma$  bytes
- Pointer to key (or object) requires  $\eta$  bytes
- For B+-tree

$$C \geq 2.\Theta.\gamma + (2.\Theta + 1).\eta$$
$$\Rightarrow \Theta_{\text{B+-tree}} = \left\lfloor \frac{C - \eta}{2(\gamma + \eta)} \right\rfloor$$

- For B-tree

# Order

- Order mainly determined by disk page size
- Disk page capacity is  $C$  bytes
- Key consumes  $\gamma$  bytes
- Pointer to key (or object) requires  $\eta$  bytes
- For B+-tree

$$C \geq 2.\Theta.\gamma + (2.\Theta + 1).\eta$$
$$\Rightarrow \Theta_{\text{B+-tree}} = \left\lfloor \frac{C - \eta}{2(\gamma + \eta)} \right\rfloor$$

- For B-tree

$$C \geq 2.\Theta.\gamma + 2.\Theta.\eta + (2.\Theta + 1).\eta$$
$$\Rightarrow \Theta_{\text{B-tree}} = \left\lfloor \frac{C - \eta}{2(\gamma + 2\eta)} \right\rfloor$$



# Example

- What is the order of a B+-tree and a B-tree with a page size of 4 KB indexing keys of 8 bytes each, and having pointers of size 4 bytes?

# Example

- What is the order of a B+-tree and a B-tree with a page size of 4 KB indexing keys of 8 bytes each, and having pointers of size 4 bytes?

- $\Theta_{\text{B+-tree}} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+4)} \right\rfloor = 170$

# Example

- What is the order of a B+-tree and a B-tree with a page size of 4 KB indexing keys of 8 bytes each, and having pointers of size 4 bytes?

- $\Theta_{B+-tree} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+4)} \right\rfloor = 170$

- $\Theta_{B-tree} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+2 \times 4)} \right\rfloor = 127$



# Example

- What is the order of a B+-tree and a B-tree with a page size of 4 KB indexing keys of 8 bytes each, and having pointers of size 4 bytes?
  - $\Theta_{B^+-tree} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+4)} \right\rfloor = 170$
  - $\Theta_{B-tree} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+2 \times 4)} \right\rfloor = 127$
- What is the height of the B+-tree and the B-tree for  $3 \times 10^7$  keys?

# Example

- What is the order of a B+-tree and a B-tree with a page size of 4 KB indexing keys of 8 bytes each, and having pointers of size 4 bytes?
  - $\Theta_{B+-tree} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+4)} \right\rfloor = 170$
  - $\Theta_{B-tree} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+2 \times 4)} \right\rfloor = 127$
- What is the height of the B+-tree and the B-tree for  $3 \times 10^7$  keys?
  - Height of B+-tree is  $\lceil \log_{2 \times 170}(3 \times 10^7) \rceil = 3$

# Example

- What is the order of a B+-tree and a B-tree with a page size of 4 KB indexing keys of 8 bytes each, and having pointers of size 4 bytes?
  - $\Theta_{B+-tree} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+4)} \right\rfloor = 170$
  - $\Theta_{B-tree} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+2 \times 4)} \right\rfloor = 127$
- What is the height of the B+-tree and the B-tree for  $3 \times 10^7$  keys?
  - Height of B+-tree is  $\lceil \log_{2 \times 170}(3 \times 10^7) \rceil = 3$
  - Height of B-tree is  $\lceil \log_{2 \times 127}(3 \times 10^7) \rceil = 4$

# Indexing Multiple Attributes

- Search keys having more than one attribute are called **composite search keys**
- Separate indices may be used
  - Union, intersection, etc. of individual results
- Multi-dimensional indexing
  - Indexing is specified by hyper-rectangles

# Indexing Multiple Attributes

- Search keys having more than one attribute are called **composite search keys**
- Separate indices may be used
  - Union, intersection, etc. of individual results
- Multi-dimensional indexing
  - Indexing is specified by hyper-rectangles
  - Space-partitioning
  - **Quadtree**, **KD-tree**, **K-d-B-tree**: Extension of BST

# Indexing Multiple Attributes

- Search keys having more than one attribute are called **composite search keys**
- **Separate indices may be used**
  - **Union, intersection, etc. of individual results**
- **Multi-dimensional indexing**
  - Indexing is specified by hyper-rectangles
  - Space-partitioning
  - **Quadtree, KD-tree, K-d-B-tree**: Extension of BST
  - Data-partitioning
  - **R-tree**: Extension of B+-tree
  - Uses **minimum bounding rectangles (MBRs)**

# Bitmap Index

- Attribute domain consists of a small number of distinct values
- A **bitmap** or a **bit vector** is an array of bits
- Each distinct value has an array of the size of the number of tuples
  - If the  $i$ -th bit is 1, tuple  $i$  has that value

Gender	Grade
Male	C
Female	A
Female	C
Male	D
Male	A

- Two sets of bit vectors
  - Male = (10011), Female = (01100)
  - A = (01001), B = (00000), C = (10100), D = (00010)

# Bitmap Operations

- Queries are answered using bitmap operations
- Example: Find the male student who got 'D'
  - $\text{Bitmap}(\text{Male}) \text{ AND } \text{Bitmap}(\text{D})$
- Null values require a special bitmap for null



# Bitmap Operations

- Queries are answered using bitmap operations
- Example: Find the male student who got 'D'
  - $\text{Bitmap}(\text{Male}) \text{ AND } \text{Bitmap}(\text{D})$
- Null values require a special bitmap for null
- O/S allows efficient bitmap operations when they are packed in word sizes

# Indices in SQL

- `create [unique] index name on rel (att)` creates an index *name* on attribute *att* of relation *rel*

```
create index idx_ctype on course (ctype);
```

# Indices in SQL

- `create [unique] index name on rel (att)` creates an index *name* on attribute *att* of relation *rel*

```
create index idx_ctype on course (ctype);
```

- Index will be used whenever attribute *att* of relation *rel* is used

# Indices in SQL

- `create [unique] index name on rel (att)` creates an index *name* on attribute *att* of relation *rel*

```
create index idx_ctype on course (ctype);
```

- Index will be used whenever attribute *att* of relation *rel* is used
- *unique* does not allow duplicates on the attributes

# Indices in SQL

- `create [unique] index name on rel (att)` creates an index *name* on attribute *att* of relation *rel*

```
create index idx_ctype on course (ctype);
```

- Index will be used whenever attribute *att* of relation *rel* is used
- *unique* does not allow duplicates on the attributes
- Often, primary keys are *implicitly* (by default) indexed

# Indices in SQL

- `create [unique] index name on rel (att)` creates an index *name* on attribute *att* of relation *rel*

```
create index idx_ctype on course (ctype);
```

- Index will be used whenever attribute *att* of relation *rel* is used
- *unique* does not allow duplicates on the attributes
- Often, primary keys are *implicitly* (by default) indexed
- Slows down modification operations as index is also modified

# Indices in SQL

- **create [unique] index name on rel (att)** creates an index *name* on attribute *att* of relation *rel*

```
create index idx_cotype on course (cotype);
```

- Index will be used whenever attribute *att* of relation *rel* is used
- *unique* does not allow duplicates on the attributes
- Often, primary keys are *implicitly* (by default) indexed
- Slows down modification operations as index is also modified
- **drop index i** deletes the index

```
drop index idx_cotype;
```