

# CS315: DATABASE SYSTEMS QUERY PROCESSING

Arnab Bhattacharya

`arnabb@cse.iitk.ac.in`

Computer Science and Engineering,  
Indian Institute of Technology, Kanpur

`http://web.cse.iitk.ac.in/~cs315/`

2<sup>nd</sup> semester, 2022-23

Tue 10:30-11:45, Thu 12:00-13:15

- Query in a high-level language is first parsed and validated

```
select cpi from students where cpi > 9;
```

# Query

- Query in a high-level language is first parsed and validated  
**select** cpi **from** students **where** cpi > 9;
- Query planner next outputs an equivalent expression in relational algebra
  - $\Pi_{cpi}(\sigma_{cpi > 9}(students))$
  - A query order tree is created

# Query

- Query in a high-level language is first parsed and validated  
**select** cpi **from** students **where** cpi > 9;
- Query planner next outputs an equivalent expression in relational algebra
  - $\Pi_{cpi}(\sigma_{cpi>9}(students))$
  - A query order tree is created
- All equivalent evaluation plans are then generated
  - $\Pi_{cpi}(\sigma_{cpi>9}(students))$
  - $\sigma_{cpi>9}(\Pi_{cpi}(students))$

# Query

- Query in a high-level language is first parsed and validated  
**select** cpi **from** students **where** cpi > 9;
- Query planner next outputs an equivalent expression in relational algebra
  - $\Pi_{cpi}(\sigma_{cpi>9}(students))$
  - A query order tree is created
- All equivalent evaluation plans are then generated
  - $\Pi_{cpi}(\sigma_{cpi>9}(students))$
  - $\sigma_{cpi>9}(\Pi_{cpi}(students))$
- Query optimizer next decides on the best evaluation plan among all equivalent plans
  - Algorithms
  - Metadata
  - Statistics of data

# Query

- Query in a high-level language is first parsed and validated  
**select** cpi **from** students **where** cpi > 9;
- Query planner next outputs an equivalent expression in relational algebra
  - $\Pi_{cpi}(\sigma_{cpi>9}(students))$
  - A query order tree is created
- All equivalent evaluation plans are then generated
  - $\Pi_{cpi}(\sigma_{cpi>9}(students))$
  - $\sigma_{cpi>9}(\Pi_{cpi}(students))$
- Query optimizer next decides on the best evaluation plan among all equivalent plans
  - Algorithms
  - Metadata
  - Statistics of data
- Query code is finally generated and processed

# Query Cost

- Factors that affect the runtime of the query
  - Disk accesses
  - CPU time
  - Network communication
  - Buffering and caching effects
- In a non-distributed setting, disk access is the most dominant factor
  - In a distributed environment, it is network communication

# Query Cost

- Factors that affect the runtime of the query
  - Disk accesses
  - CPU time
  - Network communication
  - Buffering and caching effects
- In a non-distributed setting, disk access is the most dominant factor
  - In a distributed environment, it is network communication
- Disk access can be estimated as a total of
  - Number of seeks times average seek cost
  - Number of blocks read or written times average block transfer cost
    - Write cost can be more since data is verified



# Query Cost

- Factors that affect the runtime of the query
  - Disk accesses
  - CPU time
  - Network communication
  - Buffering and caching effects
- In a non-distributed setting, disk access is the most dominant factor
  - In a distributed environment, it is network communication
- Disk access can be estimated as a total of
  - Number of seeks times average seek cost
  - Number of blocks read or written times average block transfer cost
    - Write cost can be more since data is verified
- For  $s$  seeks and  $b$  block transfers, simply estimated as  $s \times t_s + b \times t_b$ 
  - Ignores CPU time and buffer management issues

# Selection: Linear Search

- **LINEAR SEARCH**
- *Always* applicable

# Selection: Linear Search

- **LINEAR SEARCH**
- *Always* applicable
- Scan all file blocks and test each record

# Selection: Linear Search

- **LINEAR SEARCH**
- *Always* applicable
- Scan all file blocks and test each record
- Cost for a relation containing  $b$  blocks

# Selection: Linear Search

- **LINEAR SEARCH**
- *Always* applicable
- Scan all file blocks and test each record
- Cost for a relation containing  $b$  blocks
  - 1 seek
  - $b$  transfers

# Selection: Linear Search

- **LINEAR SEARCH**
- *Always* applicable
- Scan all file blocks and test each record
- Cost for a relation containing  $b$  blocks
  - 1 seek
  - $b$  transfers
- If equality on key or unique attribute, then

# Selection: Linear Search

- **LINEAR SEARCH**
- Always applicable
- Scan all file blocks and test each record
- Cost for a relation containing  $b$  blocks
  - 1 seek
  - $b$  transfers
- If equality on key or unique attribute, then
  - $b/2$  transfers on average

# Selection: Binary Search

- BINARY SEARCH
- File must be sorted on attribute which is searched



# Selection: Binary Search

- **BINARY SEARCH**
- File must be sorted on attribute which is searched
- *Equality*

# Selection: Binary Search

- **BINARY SEARCH**
- File must be sorted on attribute which is searched
- *Equality*
- Cost for locating the first record is

# Selection: Binary Search

- BINARY SEARCH
- File must be sorted on attribute which is searched
- *Equality*
- Cost for locating the first record is
  - $\lceil \log_2 b \rceil$  seeks
  - $\lceil \log_2 b \rceil$  transfers

# Selection: Binary Search

- **BINARY SEARCH**
- File must be sorted on attribute which is searched
- *Equality*
- Cost for locating the first record is
  - $\lceil \log_2 b \rceil$  seeks
  - $\lceil \log_2 b \rceil$  transfers
- If one block does not contain all records

# Selection: Binary Search

- **BINARY SEARCH**
- File must be sorted on attribute which is searched
- *Equality*
- Cost for locating the first record is
  - $\lceil \log_2 b \rceil$  seeks
  - $\lceil \log_2 b \rceil$  transfers
- If one block does not contain all records
  - Additional block transfers

# Selection: Binary Search

- **BINARY SEARCH**
- File must be sorted on attribute which is searched
- *Equality*
- Cost for locating the first record is
  - $\lceil \log_2 b \rceil$  seeks
  - $\lceil \log_2 b \rceil$  transfers
- If one block does not contain all records
  - Additional block transfers
- *Greater than*

# Selection: Binary Search

- **BINARY SEARCH**
- File must be sorted on attribute which is searched
- *Equality*
- Cost for locating the first record is
  - $\lceil \log_2 b \rceil$  seeks
  - $\lceil \log_2 b \rceil$  transfers
- If one block does not contain all records
  - Additional block transfers
- *Greater than*
  - Locate matching record
  - Traverse forward and do additional transfers

# Selection: Binary Search

- **BINARY SEARCH**
- File must be sorted on attribute which is searched
- *Equality*
- Cost for locating the first record is
  - $\lceil \log_2 b \rceil$  seeks
  - $\lceil \log_2 b \rceil$  transfers
- If one block does not contain all records
  - Additional block transfers
- *Greater than*
  - Locate matching record
  - Traverse forward and do additional transfers
- *Lesser than*



# Selection: Binary Search

- **BINARY SEARCH**
- File must be sorted on attribute which is searched
- *Equality*
- Cost for locating the first record is
  - $\lceil \log_2 b \rceil$  seeks
  - $\lceil \log_2 b \rceil$  transfers
- If one block does not contain all records
  - Additional block transfers
- *Greater than*
  - Locate matching record
  - Traverse forward and do additional transfers
- *Lesser than*
  - Scan from beginning till matching record
  - Reduces to *linear search*

# Selection: Index Search

- INDEX SEARCH
- B+-tree of height  $h$

# Selection: Index Search

- INDEX SEARCH
- B+-tree of height  $h$
- *Equality* using primary index

# Selection: Index Search

- INDEX SEARCH
- B+-tree of height  $h$
- *Equality* using primary index
  - $h + 1$  seeks, where  $h$  is the height of B+-tree
  - $h + m$  transferswhere  $m$  is the total number of blocks containing matching records
- For key,  $m = 1$

# Selection: Index Search

- INDEX SEARCH

- B+-tree of height  $h$

- *Equality* using primary index

- $h + 1$  seeks, where  $h$  is the height of B+-tree
- $h + m$  transfers

where  $m$  is the total number of blocks containing matching records

- For key,  $m = 1$

- *Equality* using secondary index

# Selection: Index Search

- INDEX SEARCH

- B+-tree of height  $h$

- Equality using primary index

- $h + 1$  seeks, where  $h$  is the height of B+-tree
- $h + m$  transfers

where  $m$  is the total number of blocks containing matching records

- For key,  $m = 1$

- Equality using secondary index

- $h + n$  seeks, where  $h$  is the height of the index tree
- $h + n$  transfers

where  $n$  is the total number of matching records, each in a separate block

- For key,  $n = 1$

## Selection: Index Search (contd.)

- *Greater than* ( $A \geq v$ ) using primary index

## Selection: Index Search (contd.)

- *Greater than* ( $A \geq v$ ) using primary index
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$
  - Scan for the rest of the relation resulting in  $n$  more transfers



## Selection: Index Search (contd.)

- *Greater than* ( $A \geq v$ ) using primary index
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$
  - Scan for the rest of the relation resulting in  $n$  more transfers
- *Lesser than* ( $A \leq v$ ) using primary index

## Selection: Index Search (contd.)

- *Greater than* ( $A \geq v$ ) using primary index
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$
  - Scan for the rest of the relation resulting in  $n$  more transfers
- *Lesser than* ( $A \leq v$ ) using primary index
  - 1 seek to the first block
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$

## Selection: Index Search (contd.)

- *Greater than* ( $A \geq v$ ) using primary index
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$
  - Scan for the rest of the relation resulting in  $n$  more transfers
- *Lesser than* ( $A \leq v$ ) using primary index
  - 1 seek to the first block
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$

or

- $n$  transfers till  $v$  is located
- Use sibling leaf pointers to locate all other matching

## Selection: Index Search (contd.)

- *Greater than* ( $A \geq v$ ) using primary index
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$
  - Scan for the rest of the relation resulting in  $n$  more transfers
- *Lesser than* ( $A \leq v$ ) using primary index
  - 1 seek to the first block
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$

or

- $n$  transfers till  $v$  is located
  - Use sibling leaf pointers to locate all other matching
- *Greater than* ( $A \geq v$ ) using secondary index

## Selection: Index Search (contd.)

- *Greater than* ( $A \geq v$ ) using primary index
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$
  - Scan for the rest of the relation resulting in  $n$  more transfers
- *Lesser than* ( $A \leq v$ ) using primary index
  - 1 seek to the first block
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$

or

- $n$  transfers till  $v$  is located
  - Use sibling leaf pointers to locate all other matching
- *Greater than* ( $A \geq v$ ) using secondary index
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$
  - Use sibling leaf pointers to locate all other matching records resulting in  $n$  more seeks and  $n$  more transfers

## Selection: Index Search (contd.)

- *Greater than ( $A \geq v$ )* using primary index
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$
  - Scan for the rest of the relation resulting in  $n$  more transfers
- *Lesser than ( $A \leq v$ )* using primary index
  - 1 seek to the first block
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$

or

- $n$  transfers till  $v$  is located
  - Use sibling leaf pointers to locate all other matching
- *Greater than ( $A \geq v$ )* using secondary index
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$
  - Use sibling leaf pointers to locate all other matching records resulting in  $n$  more seeks and  $n$  more transfers
- *Lesser than ( $A \leq v$ )* using secondary index

# Selection: Index Search (contd.)

- *Greater than ( $A \geq v$ )* using primary index
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$
  - Scan for the rest of the relation resulting in  $n$  more transfers
- *Lesser than ( $A \leq v$ )* using primary index
  - 1 seek to the first block
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$

or

- $n$  transfers till  $v$  is located
  - Use sibling leaf pointers to locate all other matching
- *Greater than ( $A \geq v$ )* using secondary index
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$
  - Use sibling leaf pointers to locate all other matching records resulting in  $n$  more seeks and  $n$  more transfers
- *Lesser than ( $A \leq v$ )* using secondary index
  - $h + 1$  seeks and  $h + 1$  transfers to locate  $v$
  - Use sibling leaf pointers to locate all other matching records resulting in  $n$  more seeks and  $n$  more transfers

# Selection on Multiple Attributes

- *Conjunction*: AND



# Selection on Multiple Attributes

- *Conjunction: AND*
  - Single index is available

# Selection on Multiple Attributes

- *Conjunction: AND*
  - Single index is available
    - For each record satisfying it, test the other attributes

# Selection on Multiple Attributes

- *Conjunction: AND*
  - Single index is available
    - For each record satisfying it, test the other attributes
  - Multiple indexes are available

# Selection on Multiple Attributes

- *Conjunction: AND*
  - Single index is available
    - For each record satisfying it, test the other attributes
  - Multiple indexes are available
    - Use one likely to produce lesser tuples (e.g., equality)

# Selection on Multiple Attributes

- *Conjunction: AND*
  - Single index is available
    - For each record satisfying it, test the other attributes
  - Multiple indexes are available
    - Use one likely to produce lesser tuples (e.g., equality)
  - Intersection
    - If no attribute has an index, test explicitly

# Selection on Multiple Attributes

- *Conjunction: AND*
  - Single index is available
    - For each record satisfying it, test the other attributes
  - Multiple indexes are available
    - Use one likely to produce lesser tuples (e.g., equality)
  - Intersection
    - If no attribute has an index, test explicitly
- *Disjunction: OR*

# Selection on Multiple Attributes

- *Conjunction: AND*
  - Single index is available
    - For each record satisfying it, test the other attributes
  - Multiple indexes are available
    - Use one likely to produce lesser tuples (e.g., equality)
  - Intersection
    - If no attribute has an index, test explicitly
- *Disjunction: OR*
  - Union
    - If some attribute does not have an index, then linear scan

# Selection on Multiple Attributes

- *Conjunction: AND*
  - Single index is available
    - For each record satisfying it, test the other attributes
  - Multiple indexes are available
    - Use one likely to produce lesser tuples (e.g., equality)
  - Intersection
    - If no attribute has an index, test explicitly
- *Disjunction: OR*
  - Union
    - If some attribute does not have an index, then linear scan
  - Negation of conjunction
    - May be very inefficient



# Selection on Multiple Attributes

- *Conjunction: AND*
  - Single index is available
    - For each record satisfying it, test the other attributes
  - Multiple indexes are available
    - Use one likely to produce lesser tuples (e.g., equality)
  - Intersection
    - If no attribute has an index, test explicitly
- *Disjunction: OR*
  - Union
    - If some attribute does not have an index, then linear scan
  - Negation of conjunction
    - May be very inefficient
- *Negation of equality ( $A \neq v$ )*


# Selection on Multiple Attributes

- *Conjunction: AND*
  - Single index is available
    - For each record satisfying it, test the other attributes
  - Multiple indexes are available
    - Use one likely to produce lesser tuples (e.g., equality)
  - Intersection
    - If no attribute has an index, test explicitly
- *Disjunction: OR*
  - Union
    - If some attribute does not have an index, then linear scan
  - Negation of conjunction
    - May be very inefficient
- *Negation of equality ( $A \neq v$ )*
  - Linear scan
  - Index selects leaf pointers for which corresponding records will not be retrieved

# Selection on Multiple Attributes

- *Conjunction: AND*
  - Single index is available
    - For each record satisfying it, test the other attributes
  - Multiple indexes are available
    - Use one likely to produce lesser tuples (e.g., equality)
  - Intersection
    - If no attribute has an index, test explicitly
- *Disjunction: OR*
  - Union
    - If some attribute does not have an index, then linear scan
  - Negation of conjunction
    - May be very inefficient
- *Negation of equality ( $A \neq v$ )*
  - Linear scan
  - Index selects leaf pointers for which corresponding records will not be retrieved
- Negation of comparison

# Selection on Multiple Attributes

- *Conjunction: AND*
  - Single index is available
    - For each record satisfying it, test the other attributes
  - Multiple indexes are available 
    - Use one likely to produce less tuples (e.g., equality)
  - Intersection
    - If no attribute has an index, test explicitly
- *Disjunction: OR*
  - Union
    - If some attribute does not have an index, then linear scan
  - Negation of conjunction
    - May be very inefficient
- *Negation of equality ( $A \neq v$ )*
  - Linear scan
  - Index selects leaf pointers for which corresponding records will not be retrieved
- Negation of comparison is just another comparison

# Sorting

- Display purposes
- Certain operations such as join can be implemented efficiently on sorted relations

# Sorting

- Display purposes
- Certain operations such as join can be implemented efficiently on sorted relations
- Index provides a *logical* sorted view

# Sorting

- Display purposes
- Certain operations such as join can be implemented efficiently on sorted relations
- Index provides a *logical* sorted view
- Tuples need to be *physically* sorted

# Sorting

- Display purposes
- Certain operations such as join can be implemented efficiently on sorted relations
- Index provides a *logical* sorted view
- Tuples need to be *physically* sorted
- When the relation fits in memory, **QUICKSORT** can be used
- When it does not, **external sorting** algorithms are used



# Sorting

- Display purposes
- Certain operations such as join can be implemented efficiently on sorted relations
- Index provides a *logical* sorted view
- Tuples need to be *physically* sorted
- When the relation fits in memory, **QUICKSORT** can be used
- When it does not, **external sorting** algorithms are used
- **EXTERNAL MERGESORT** or **EXTERNAL SORT-MERGE**

# External Mergesort

- Assume only  $m$  blocks can be put into memory
- Size of relation is more than  $m$  blocks

# External Mergesort

- Assume only  $m$  blocks can be put into memory
- Size of relation is more than  $m$  blocks
- Create sorted **runs**
  - Read  $m$  blocks at a time
  - Sort them in-memory using any algorithm such as quicksort
  - Write them back to disk

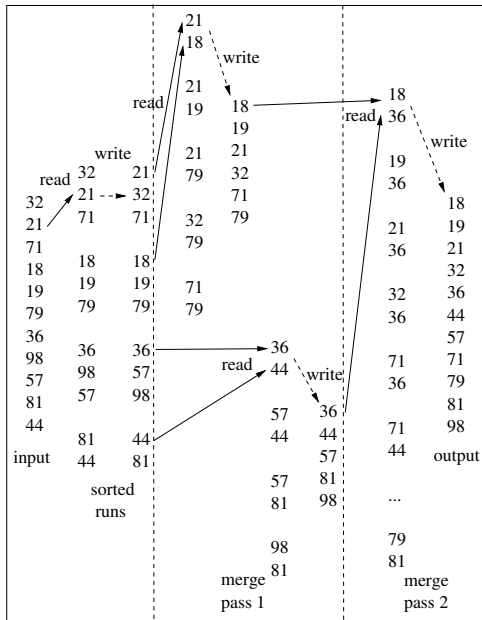
# External Mergesort

- Assume only  $m$  blocks can be put into memory
- Size of relation is more than  $m$  blocks
- Create sorted **runs**
  - Read  $m$  blocks at a time
  - Sort them in-memory using any algorithm such as quicksort
  - Write them back to disk
- Merge  $m - 1$  runs ( **$(m - 1)$ -way merge**)
  - Read in first block of  $m - 1$  runs
  - Output the first record to *buffer* block ( $m$ -th block in memory)
  - Continue till buffer block is full
  - Write buffer block to disk
  - When a block of a particular run is exhausted, read in the next block of the run

# External Mergesort

- Assume only  $m$  blocks can be put into memory
- Size of relation is more than  $m$  blocks
- Create sorted runs
  - Read  $m$  blocks at a time
  - Sort them in-memory using any algorithm such as quicksort
  - Write them back to disk
- Merge  $m - 1$  runs ( $(m - 1)$ -way merge)
  - Read in first block of  $m - 1$  runs
  - Output the first record to *buffer* block ( $m$ -th block in memory)
  - Continue till buffer block is full
  - Write buffer block to disk
  - When a block of a particular run is exhausted, read in the next block of the run
- Continue with  $(m - 1)$ -way merge till the number of sorted runs is less than  $m$
- The last  $(m - 1)$ -way merge sorts the relation

# Example



# Cost of External Mergesort

- Total number of blocks is  $b$
- Initial number of sorted runs is

# Cost of External Mergesort

- Total number of blocks is  $b$
- Initial number of sorted runs is  $n = \lceil b/m \rceil$
- In each merge pass,  $m - 1$  runs are sorted
- Therefore, total number of merge passes required is



# Cost of External Mergesort

- Total number of blocks is  $b$
- Initial number of sorted runs is  $n = \lceil b/m \rceil$
- In each merge pass,  $m - 1$  runs are sorted
- Therefore, total number of merge passes required is  $r = \lceil \log_{m-1} n \rceil$
- During each of these passes and the first pass, all blocks are read and written
- There is an initial pass of creating runs
- Hence, total number of block transfers is

# Cost of External Mergesort

- Total number of blocks is  $b$
- Initial number of sorted runs is  $n = \lceil b/m \rceil$
- In each merge pass,  $m - 1$  runs are sorted
- Therefore, total number of merge passes required is  $r = \lceil \log_{m-1} n \rceil$
- During each of these passes and the first pass, all blocks are read and written
- There is an initial pass of creating runs
- Hence, total number of block transfers is  $2br + 2b$

## Cost of External Mergesort (contd.)

- Initial pass to create sorted runs reads  $m$  blocks at a time
- Therefore, number of seeks is  $2n$  for reading and writing

## Cost of External Mergesort (contd.)

- Initial pass to create sorted runs reads  $m$  blocks at a time
- Therefore, number of seeks is  $2n$  for reading and writing
- During the merge passes, blocks from different runs may not be read consecutively
- Consequently, each read and write for another run may move the disk head away, thereby requiring a seek every time
- Hence, number of seeks for these passes is  $2br$
- Therefore, total number of seeks is

## Cost of External Mergesort (contd.)

- Initial pass to create sorted runs reads  $m$  blocks at a time
- Therefore, number of seeks is  $2n$  for reading and writing
- During the merge passes, blocks from different runs may not be read consecutively
- Consequently, each read and write for another run may move the disk head away, thereby requiring a seek every time
- Hence, number of seeks for these passes is  $2br$
- Therefore, total number of seeks is  $2n + 2br$  in the worst case

- Different join algorithms
  - NESTED-LOOP JOIN
  - BLOCK NESTED-LOOP JOIN
  - INDEXED NESTED-LOOP JOIN
  - MERGE JOIN
  - HASH JOIN
- Choice depends on cost estimates

# Nested-Loop Join

- Applicable for any kind of join
- For each record  $t_r \in r$  and for each record  $t_s \in s$ , if  $t_r \bowtie t_s$  satisfies the join condition, add it to result
- **Outer relation**  $r$ : outer loop; **inner relation**  $s$ : inner loop



# Block Nested-Loop Join

- Applicable for any kind of join
- *Disk block* aware version of nested-loop
- For each block  $l_r \in r$  and for each block  $l_s \in s$ , test if every record  $t_r \in l_r$  and  $t_s \in l_s$  satisfies the join condition; if so, add to the result



# Cost of Block Nested-Loop Join: Minimal

- Minimal assumption that only 2 blocks fit in memory

# Cost of Block Nested-Loop Join: Minimal

- Minimal assumption that only 2 blocks fit in memory
- Block transfers
  - $b_s$  transfers every time a block  $l_r$  is read
  - $b_r$  transfers for blocks in  $r$
  - Therefore, total is  $b_r + b_r \times b_s$  transfers

# Cost of Block Nested-Loop Join: Minimal

- Minimal assumption that only 2 blocks fit in memory
- Block transfers
  - $b_s$  transfers every time a block  $l_r$  is read
  - $b_r$  transfers for blocks in  $r$
  - Therefore, total is  $b_r + b_r \times b_s$  transfers
- Seeks
  - 1 seek for records in  $s$  every time a block  $l_r$  is read
  - $b_r$  seeks for records in  $r$
  - Therefore, total is  $b_r + b_r = 2b_r$  seeks
- Smaller relation should be

# Cost of Block Nested-Loop Join: Minimal

- Minimal assumption that only 2 blocks fit in memory
- Block transfers
  - $b_s$  transfers every time a block  $l_r$  is read
  - $b_r$  transfers for blocks in  $r$
  - Therefore, total is  $b_r + b_r \times b_s$  transfers
- Seeks
  - 1 seek for records in  $s$  every time a block  $l_r$  is read
  - $b_r$  seeks for records in  $r$
  - Therefore, total is  $b_r + b_r = 2b_r$  seeks
- Smaller relation should be *outer*

# Cost of Block Nested-Loop Join

- Assume  $m$  blocks of memory to be available

# Cost of Block Nested-Loop Join

- Assume  $m$  blocks of memory to be available
- 1 left for inner relation  $s$
- $m - 1$  blocks read at a time from outer relation  $r$

# Cost of Block Nested-Loop Join

- Assume  $m$  blocks of memory to be available
- 1 left for inner relation  $s$
- $m - 1$  blocks read at a time from outer relation  $r$
- Number of such runs is  $n = \lceil b_r / (m - 1) \rceil$

# Cost of Block Nested-Loop Join

- Assume  $m$  blocks of memory to be available
- 1 left for inner relation  $s$
- $m - 1$  blocks read at a time from outer relation  $r$
- Number of such runs is  $n = \lceil b_r / (m - 1) \rceil$
- Block transfers
  - For  $r$ :  $b_r$
  - For  $s$ :  $n \times b_s$



# Cost of Block Nested-Loop Join

- Assume  $m$  blocks of memory to be available
- 1 left for inner relation  $s$
- $m - 1$  blocks read at a time from outer relation  $r$
- Number of such runs is  $n = \lceil b_r / (m - 1) \rceil$
- Block transfers
  - For  $r$ :  $b_r$
  - For  $s$ :  $n \times b_s$
- Seeks
  - For  $r$ :  $n$
  - For  $s$ :  $n$
- Smaller relation should be

# Cost of Block Nested-Loop Join

- Assume  $m$  blocks of memory to be available
- 1 left for inner relation  $s$
- $m - 1$  blocks read at a time from outer relation  $r$
- Number of such runs is  $n = \lceil b_r / (m - 1) \rceil$
- Block transfers
  - For  $r$ :  $b_r$
  - For  $s$ :  $n \times b_s$
- Seeks
  - For  $r$ :  $n$
  - For  $s$ :  $n$
- Smaller relation should be *outer*



# Example

- $r$ : 40000 tuples at 200 per block:

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block:

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer (+1 for inner)

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer (+1 for inner)
- Outer should be

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer (+1 for inner)
- Outer should be  $r$



# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer (+1 for inner)
- Outer should be  $r$
- Number of runs is

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer (+1 for inner)
- Outer should be  $r$
- Number of runs is  $n = \lceil 200/25 \rceil = 8$

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer (+1 for inner)
- Outer should be  $r$
- Number of runs is  $n = \lceil 200/25 \rceil = 8$
- In each run,  $s$  requires 1 seek and 250 transfers
- Total is 8 seeks and 2000 transfers

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer (+1 for inner)
- Outer should be  $r$
- Number of runs is  $n = \lceil 200/25 \rceil = 8$
- In each run,  $s$  requires 1 seek and 250 transfers
- Total is 8 seeks and 2000 transfers
- $r$  requires 8 seeks and 200 transfers

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer (+1 for inner)
- Outer should be  $r$
- Number of runs is  $n = \lceil 200/25 \rceil = 8$
- In each run,  $s$  requires 1 seek and 250 transfers
- Total is 8 seeks and 2000 transfers
- $r$  requires 8 seeks and 200 transfers
- Total is 16 seeks and 2200 transfers

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer (+1 for inner)
- Outer should be  $r$
- Number of runs is  $n = \lceil 200/25 \rceil = 8$
- In each run,  $s$  requires 1 seek and 250 transfers
- Total is 8 seeks and 2000 transfers
- $r$  requires 8 seeks and 200 transfers
- Total is 16 seeks and 2200 transfers
- If  $s$  made outer,

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer (+1 for inner)
- Outer should be  $r$
- Number of runs is  $n = \lceil 200/25 \rceil = 8$
- In each run,  $s$  requires 1 seek and 250 transfers
- Total is 8 seeks and 2000 transfers
- $r$  requires 8 seeks and 200 transfers
- Total is 16 seeks and 2200 transfers
- If  $s$  made outer, 20 seeks and 2250 transfers

# Indexed Nested-Loop Join

- Indexed version of the block nested-loop algorithm
- Applicable when inner relation has an index on the joining attribute
- For each block  $l_r \in r$  and for each record  $t_r \in l_r$ , use index on  $s$  to locate records  $t_s \in s$  that satisfies the join condition



# Indexed Nested-Loop Join

- Indexed version of the block nested-loop algorithm
- Applicable when inner relation has an index on the joining attribute
- For each block  $I_r \in r$  and for each record  $t_r \in I_r$ , use index on  $s$  to locate records  $t_s \in s$  that satisfies the join condition
- Most effective when the join condition is equality

# Cost of Indexed Nested-Loop Join

- Assumption is that  $m$  blocks are available in memory

# Cost of Indexed Nested-Loop Join

- Assumption is that  $m$  blocks are available in memory
- $m - 1$  blocks from outer relation  $r$  can be read at a time

# Cost of Indexed Nested-Loop Join

- Assumption is that  $m$  blocks are available in memory
- $m - 1$  blocks from outer relation  $r$  can be read at a time
- Assume cost of searching index on  $s$  is  $c_s$  seeks and  $c_t$  transfers
  - Typically,  $c_s$  and  $c_t$  are height  $h$  of B+-tree

# Cost of Indexed Nested-Loop Join

- Assumption is that  $m$  blocks are available in memory
- $m - 1$  blocks from outer relation  $r$  can be read at a time
- Assume cost of searching index on  $s$  is  $c_s$  seeks and  $c_t$  transfers
  - Typically,  $c_s$  and  $c_t$  are height  $h$  of B+-tree
- Cost for  $r$  is  $n = \lceil b_r / (m - 1) \rceil$  seeks and  $b_r$  transfers
- Therefore, total is  $n + n_r \times c_s$  seeks and  $b_r + n_r \times c_t$  transfers
  - Block processing done to reduce  $n_r$  to  $b_r$



# Cost of Indexed Nested-Loop Join

- Assumption is that  $m$  blocks are available in memory
- $m - 1$  blocks from outer relation  $r$  can be read at a time
- Assume cost of searching index on  $s$  is  $c_s$  seeks and  $c_t$  transfers
  - Typically,  $c_s$  and  $c_t$  are height  $h$  of B+-tree
- Cost for  $r$  is  $n = \lceil b_r / (m - 1) \rceil$  seeks and  $b_r$  transfers
- Therefore, total is  $n + n_r \times c_s$  seeks and  $b_r + n_r \times c_t$  transfers
  - Block processing done to reduce  $n_r$  to  $b_r$
- If index is available on both relations, like block nested-loop, smaller relation should be

# Cost of Indexed Nested-Loop Join

- Assumption is that  $m$  blocks are available in memory
- $m - 1$  blocks from outer relation  $r$  can be read at a time
- Assume cost of searching index on  $s$  is  $c_s$  seeks and  $c_t$  transfers
  - Typically,  $c_s$  and  $c_t$  are height  $h$  of B+-tree
- Cost for  $r$  is  $n = \lceil b_r / (m - 1) \rceil$  seeks and  $b_r$  transfers
- Therefore, total is  $n + n_r \times c_s$  seeks and  $b_r + n_r \times c_t$  transfers
  - Block processing done to reduce  $n_r$  to  $b_r$
- If index is available on both relations, like block nested-loop, smaller relation should be outer

# Cost of Indexed Nested-Loop Join

- Assumption is that  $m$  blocks are available in memory
- $m - 1$  blocks from outer relation  $r$  can be read at a time
- Assume cost of searching index on  $s$  is  $c_s$  seeks and  $c_t$  transfers
  - Typically,  $c_s$  and  $c_t$  are height  $h$  of B+-tree
- Cost for  $r$  is  $n = \lceil b_r / (m - 1) \rceil$  seeks and  $b_r$  transfers 
- Therefore, total is  $n + n_r \times c_s$  seeks and  $b_r + n_r \times c_t$  transfers
  - Block processing done to reduce  $n_r$  to  $b_r$
- If index is available on both relations, like block nested-loop, smaller relation should be outer
- Generally, all levels of B+-tree are held in memory except the last 
  - Then, cost of index search falls to  $c_s = c_t = 1$



# Example

- $r$ : 40000 tuples at 200 per block:

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block:

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer and enough for inner

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer and enough for inner
- Outer should be

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer and enough for inner
- Outer should be  $s$

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer and enough for inner
- Outer should be  $s$
- Number of runs is

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer and enough for inner
- Outer should be  $s$
- Number of runs is  $n = \lceil 250/25 \rceil = 10$



# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer and enough for inner
- Outer should be  $s$
- Number of runs is  $n = \lceil 250/25 \rceil = 10$
- In each run,  $r$  requires 1 seek and 1 transfer
- Total is 25000 seeks and 25000 transfers

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer and enough for inner
- Outer should be  $s$
- Number of runs is  $n = \lceil 250/25 \rceil = 10$
- In each run,  $r$  requires 1 seek and 1 transfer
- Total is 25000 seeks and 25000 transfers
- $s$  requires 10 seeks and 250 transfers

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 25$  for outer and enough for inner
- Outer should be  $s$
- Number of runs is  $n = \lceil 250/25 \rceil = 10$
- In each run,  $r$  requires 1 seek and 1 transfer
- Total is 25000 seeks and 25000 transfers
- $s$  requires 10 seeks and 250 transfers
- Total is 25010 seeks and 25250 transfers

# Merge Join or Sort-Merge Join

- Applicable only when the join condition is *equality*
- Relations need to be sorted according to the joining attribute

# Merge Join or Sort-Merge Join

- Applicable only when the join condition is *equality*
- Relations need to be sorted according to the joining attribute
- Proceed in sorted order on two relations
- If records match, output; otherwise, advance to next record
- Join step is similar to merge step in mergesort

# Merge Join or Sort-Merge Join

- Applicable only when the join condition is *equality*
- Relations need to be sorted according to the joining attribute
- Proceed in sorted order on two relations
- If records match, output; otherwise, advance to next record
- Join step is similar to merge step in mergesort
- If relations are not sorted, secondary index on attributes can be used
- **HYBRID MERGE JOIN** algorithm merges sorted records in one relation with B+-tree leaves of other relation

# Cost of Merge Join

- Each record is read only once
- Consequently, each block is read only once

# Cost of Merge Join

- Each record is read only once
- Consequently, each block is read only once
- However, blocks may be read in an interleaved manner



# Cost of Merge Join

- Each record is read only once
- Consequently, each block is read only once
- However, blocks may be read in an interleaved manner
- If only 2 blocks of memory are available (1 for each), cost is

# Cost of Merge Join

- Each record is read only once
- Consequently, each block is read only once
- However, blocks may be read in an interleaved manner
- If only 2 blocks of memory are available (1 for each), cost is
  - $b_r + b_s$  transfers
  - $b_r + b_s$  seeks

# Cost of Merge Join

- Each record is read only once
- Consequently, each block is read only once
- However, blocks may be read in an interleaved manner
- If only 2 blocks of memory are available (1 for each), cost is
  - $b_r + b_s$  transfers
  - $b_r + b_s$  seeks
- If  $m = m'/2$  blocks of memory are available for each relation, cost is

# Cost of Merge Join

- Each record is read only once
- Consequently, each block is read only once
- However, blocks may be read in an interleaved manner
- If only 2 blocks of memory are available (1 for each), cost is
  - $b_r + b_s$  transfers
  - $b_r + b_s$  seeks
- If  $m = m'/2$  blocks of memory are available for each relation, cost is  $\lceil b_r/m \rceil + \lceil b_s/m \rceil$  seeks and  $b_r + b_s$  transfers

# Example

- $r$ : 40000 tuples at 200 per block:

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block:

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 13$  for each



# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 13$  for each
- Number of seeks for  $r$  is

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 13$  for each
- Number of seeks for  $r$  is  $\lceil 250/13 \rceil = 20$
- Number of seeks for  $s$  is

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 13$  for each
- Number of seeks for  $r$  is  $\lceil 250/13 \rceil = 20$
- Number of seeks for  $s$  is  $\lceil 200/13 \rceil = 16$

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 13$  for each
- Number of seeks for  $r$  is  $\lceil 250/13 \rceil = 20$
- Number of seeks for  $s$  is  $\lceil 200/13 \rceil = 16$
- Total is

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is  $m = 13$  for each
- Number of seeks for  $r$  is  $\lceil 250/13 \rceil = 20$
- Number of seeks for  $s$  is  $\lceil 200/13 \rceil = 16$
- Total is 36 seeks and 450 transfers

# Hash Join

- Applicable only when the join condition is *equality*
- If record  $t_r$  and  $t_s$  match, they must hash to same value, and thus, only partitions with the same hash value need to be compared

# Hash Join

- Applicable only when the join condition is *equality*
- If record  $t_r$  and  $t_s$  match, they must hash to same value, and thus, only partitions with the same hash value need to be compared
- **PARTITION HASH JOIN**
- Hash function  $h$  to partition records of *both* relations into  $n$  partitions:  $h : t \in (r \cup s) \rightarrow \{0, \dots, n-1\}$

# Hash Join

- Applicable only when the join condition is *equality*
- If record  $t_r$  and  $t_s$  match, they must hash to same value, and thus, only partitions with the same hash value need to be compared
- **PARTITION HASH JOIN**
- Hash function  $h$  to partition records of *both* relations into  $n$  partitions:  $h : t \in (r \cup s) \rightarrow \{0, \dots, n-1\}$
- **Partitioning** (or **building**) phase partitions the relations and writes them to disk



# Hash Join

- Applicable only when the join condition is *equality*
- If record  $t_r$  and  $t_s$  match, they must hash to same value, and thus, only partitions with the same hash value need to be compared
- **PARTITION HASH JOIN**
- Hash function  $h$  to partition records of *both* relations into  $n$  partitions:  $h : t \in (r \cup s) \rightarrow \{0, \dots, n-1\}$
- **Partitioning** (or **building**) phase partitions the relations and writes them to disk
- **Matching** (or **probing**) phase
  - Partition  $P_i(s)$  is read into memory:  $t_s \in P_i(s) \iff h(t_s) = i$
  - A finer-grained in-memory hash may be constructed on  $P_i(s)$

# Hash Join

- Applicable only when the join condition is *equality*
- If record  $t_r$  and  $t_s$  match, they must hash to same value, and thus, only partitions with the same hash value need to be compared
- **PARTITION HASH JOIN**
- Hash function  $h$  to partition records of *both* relations into  $n$  partitions:  $h : t \in (r \cup s) \rightarrow \{0, \dots, n-1\}$
- **Partitioning** (or **building**) phase partitions the relations and writes them to disk
- **Matching** (or **probing**) phase
  - Partition  $P_i(s)$  is read into memory:  $t_s \in P_i(s) \iff h(t_s) = i$
  - A finer-grained in-memory hash may be constructed on  $P_i(s)$
  - Partition  $P_i(r)$  is read block by block:  $t_r \in P_i(r) \iff h(t_r) = i$
  - For each tuple in  $P_i(r)$ , in-memory hash of  $P_i(s)$  is searched

# Hash Join

- Applicable only when the join condition is *equality*
- If record  $t_r$  and  $t_s$  match, they must hash to same value, and thus, only partitions with the same hash value need to be compared
- **PARTITION HASH JOIN**
- Hash function  $h$  to partition records of *both* relations into  $n$  partitions:  $h : t \in (r \cup s) \rightarrow \{0, \dots, n-1\}$
- **Partitioning** (or **building**) phase partitions the relations and writes them to disk
- **Matching** (or **probing**) phase
  - Partition  $P_i(s)$  is read into memory:  $t_s \in P_i(s) \iff h(t_s) = i$
  - A finer-grained in-memory hash may be constructed on  $P_i(s)$
  - Partition  $P_i(r)$  is read block by block:  $t_r \in P_i(r) \iff h(t_r) = i$
  - For each tuple in  $P_i(r)$ , in-memory hash of  $P_i(s)$  is searched
  - After partition  $i$  is done, new partitions are loaded

# Hash Join

- Applicable only when the join condition is *equality*
- If record  $t_r$  and  $t_s$  match, they must hash to same value, and thus, only partitions with the same hash value need to be compared
- **PARTITION HASH JOIN**
- Hash function  $h$  to partition records of *both* relations into  $n$  partitions:  $h : t \in (r \cup s) \rightarrow \{0, \dots, n-1\}$
- **Partitioning** (or **building**) phase partitions the relations and writes them to disk
- **Matching** (or **probing**) phase
  - Partition  $P_i(s)$  is read into memory:  $t_s \in P_i(s) \iff h(t_s) = i$
  - A finer-grained in-memory hash may be constructed on  $P_i(s)$
  - Partition  $P_i(r)$  is read block by block:  $t_r \in P_i(r) \iff h(t_r) = i$
  - For each tuple in  $P_i(r)$ , in-memory hash of  $P_i(s)$  is searched
  - After partition  $i$  is done, new partitions are loaded
- $s$  is called **build input**
- $r$  is called **probe input**

# Hash Join

- Applicable only when the join condition is *equality*
- If record  $t_r$  and  $t_s$  match, they must hash to same value, and thus, only partitions with the same hash value need to be compared
- **PARTITION HASH JOIN**
- Hash function  $h$  to partition records of *both* relations into  $n$  partitions:  $h : t \in (r \cup s) \rightarrow \{0, \dots, n-1\}$
- **Partitioning** (or **building**) phase partitions the relations and writes them to disk
- **Matching** (or **probing**) phase
  - Partition  $P_i(s)$  is read into memory:  $t_s \in P_i(s) \iff h(t_s) = i$
  - A finer-grained in-memory hash may be constructed on  $P_i(s)$
  - Partition  $P_i(r)$  is read block by block:  $t_r \in P_i(r) \iff h(t_r) = i$
  - For each tuple in  $P_i(r)$ , in-memory hash of  $P_i(s)$  is searched
  - After partition  $i$  is done, new partitions are loaded
- $s$  is called **build input**
- $r$  is called **probe input**
- Each partition of build relation must be stored in memory
- Smaller relation should be

# Hash Join

- Applicable only when the join condition is *equality*
- If record  $t_r$  and  $t_s$  match, they must hash to same value, and thus, only partitions with the same hash value need to be compared
- **PARTITION HASH JOIN**
- Hash function  $h$  to partition records of *both* relations into  $n$  partitions:  $h : t \in (r \cup s) \rightarrow \{0, \dots, n-1\}$
- **Partitioning** (or **building**) phase partitions the relations and writes them to disk
- **Matching** (or **probing**) phase
  - Partition  $P_i(s)$  is read into memory:  $t_s \in P_i(s) \iff h(t_s) = i$
  - A finer-grained in-memory hash may be constructed on  $P_i(s)$
  - Partition  $P_i(r)$  is read block by block:  $t_r \in P_i(r) \iff h(t_r) = i$
  - For each tuple in  $P_i(r)$ , in-memory hash of  $P_i(s)$  is searched
  - After partition  $i$  is done, new partitions are loaded
- $s$  is called **build input**
- $r$  is called **probe input**
- Each partition of build relation must be stored in memory
- Smaller relation should be *build*

# Single Pass

- Each of the  $n$  partitions of build input  $s$  should fit into memory
- If build relation has  $b_s$  blocks, each partition has roughly  $b_s/n$  blocks

# Single Pass

- Each of the  $n$  partitions of build input  $s$  should fit into memory
- If build relation has  $b_s$  blocks, each partition has roughly  $b_s/n$  blocks
- Assume an available memory of  $m$  blocks for build



# Single Pass

- Each of the  $n$  partitions of build input  $s$  should fit into memory
- If build relation has  $b_s$  blocks, each partition has roughly  $b_s/n$  blocks
- Assume an available memory of  $m$  blocks for build
- If each partition fits, then  $m > b_s/n$  or  $n > b_s/m$

# Single Pass

- Each of the  $n$  partitions of build input  $s$  should fit into memory
- If build relation has  $b_s$  blocks, each partition has roughly  $b_s/n$  blocks
- Assume an available memory of  $m$  blocks for build
- If each partition fits, then  $m > b_s/n$  or  $n > b_s/m$
- During partitioning, at least 1 block of each partition should fit in memory
- Thus,  $m > n$

# Single Pass

- Each of the  $n$  partitions of build input  $s$  should fit into memory
- If build relation has  $b_s$  blocks, each partition has roughly  $b_s/n$  blocks
- Assume an available memory of  $m$  blocks for build
- If each partition fits, then  $m > b_s/n$  or  $n > b_s/m$
- During partitioning, at least 1 block of each partition should fit in memory
- Thus,  $m > n$
- Combining,

# Single Pass

- Each of the  $n$  partitions of build input  $s$  should fit into memory
- If build relation has  $b_s$  blocks, each partition has roughly  $b_s/n$  blocks
- Assume an available memory of  $m$  blocks for build
- If each partition fits, then  $m \geq b_s/n$  or  $n \geq b_s/m$
- During partitioning, at least 1 block of each partition should fit in memory
- Thus,  $m \geq n$
- Combining,  $m \geq n \geq b_s/m$  or,  $m \geq \sqrt{b_s}$
- If not so, then partitioning *cannot* be done in one pass

# Recursive Partitioning

- If  $m < \sqrt{b_s}$ , then partitioning *cannot* be done in one pass

# Recursive Partitioning

- If  $m < \sqrt{b_s}$ , then partitioning *cannot* be done in one pass
- **Recursive partitioning** is employed

# Recursive Partitioning

- If  $m < \sqrt{b_s}$ , then partitioning *cannot* be done in one pass
- **Recursive partitioning** is employed
- Initially, number of partitions  $n_s$  is chosen to be at most  $m$
- However, each partition now does *not* fit into memory
- Thus, each partition is read and re-partitioned till each smaller partition fits into memory

# Recursive Partitioning

- If  $m < \sqrt{b_s}$ , then partitioning *cannot* be done in one pass
- **Recursive partitioning** is employed
- Initially, number of partitions  $n_s$  is chosen to be at most  $m$
- However, each partition now does *not* fit into memory
- Thus, each partition is read and re-partitioned till each smaller partition fits into memory
- Number of passes is



# Recursive Partitioning

- If  $m < \sqrt{b_s}$ , then partitioning *cannot* be done in one pass
- **Recursive partitioning** is employed
- Initially, number of partitions  $n_s$  is chosen to be at most  $m$
- However, each partition now does *not* fit into memory
- Thus, each partition is read and re-partitioned till each smaller partition fits into memory
- Number of passes is  $p = \lfloor \log_m b_s \rfloor$

# Recursive Partitioning

- If  $m < \sqrt{b_s}$ , then partitioning *cannot* be done in one pass
- **Recursive partitioning** is employed
- Initially, number of partitions  $n_s$  is chosen to be at most  $m$
- However, each partition now does *not* fit into memory
- Thus, each partition is read and re-partitioned till each smaller partition fits into memory
- Number of passes is  $p = \lfloor \log_m b_s \rfloor$
- **HYBRID HASH JOIN** when more memory is available

# Recursive Partitioning

- If  $m < \sqrt{b_s}$ , then partitioning *cannot* be done in one pass
- **Recursive partitioning** is employed
- Initially, number of partitions  $n_s$  is chosen to be at most  $m$
- However, each partition now does *not* fit into memory
- Thus, each partition is read and re-partitioned till each smaller partition fits into memory
- **Number of passes** is  $p = \lfloor \log_m b_s \rfloor$
- **HYBRID HASH JOIN** when more memory is available
- Entire build relation  $s$  can be in memory
- Retain the first partition of build relation in memory

# Example

- Suppose, build input has  $b_s = 120$  blocks

# Example

- Suppose, build input has  $b_s = 120$  blocks
- Available memory is  $m = 12$  blocks

# Example

- Suppose, build input has  $b_s = 120$  blocks
- Available memory is  $m = 12$  blocks
  - Partitioning produces  $120/12 = 10$  blocks that fit
  - Thus, only 1 pass is required

# Example

- Suppose, build input has  $b_s = 120$  blocks
- Available memory is  $m = 12$  blocks
  - Partitioning produces  $120/12 = 10$  blocks that fit
  - Thus, only 1 pass is required
- Available memory is  $m = 10$  blocks

# Example

- Suppose, build input has  $b_s = 120$  blocks
- Available memory is  $m = 12$  blocks
  - Partitioning produces  $120/12 = 10$  blocks that fit
  - Thus, only 1 pass is required
- Available memory is  $m = 10$  blocks
  - Partitioning produces  $120/10 = 12$  blocks that do not fit
  - Thus, 12 blocks are partitioned again to  $12/10 = 2$  blocks each
  - This now fits
  - Thus, 2 passes are required



# Example

- Suppose, build input has  $b_s = 120$  blocks
- Available memory is  $m = 12$  blocks
  - Partitioning produces  $120/12 = 10$  blocks that fit
  - Thus, only 1 pass is required
- Available memory is  $m = 10$  blocks
  - Partitioning produces  $120/10 = 12$  blocks that do not fit
  - Thus, 12 blocks are partitioned again to  $12/10 = 2$  blocks each
  - This now fits
  - Thus, 2 passes are required
- Available memory is  $m = 4$  blocks

# Example

- Suppose, build input has  $b_s = 120$  blocks
- Available memory is  $m = 12$  blocks
  - Partitioning produces  $120/12 = 10$  blocks that fit
  - Thus, only 1 pass is required
- Available memory is  $m = 10$  blocks
  - Partitioning produces  $120/10 = 12$  blocks that do not fit
  - Thus, 12 blocks are partitioned again to  $12/10 = 2$  blocks each
  - This now fits
  - Thus, 2 passes are required
- Available memory is  $m = 4$  blocks
  - Partitioning produces  $120/4 = 30$  and then  $30/4 = 8$  and then  $8/4 = 2$  blocks
  - Thus, 3 passes are required

# Cost of Hash Join

- Initial reading of a relation  $i$  requires  $b_i$  transfers

# Cost of Hash Join

- Initial reading of a relation  $i$  requires  $b_i$  transfers
- Total number of blocks after partitioning is at most

# Cost of Hash Join

- Initial reading of a relation  $i$  requires  $b_i$  transfers
- Total number of blocks after partitioning is at most  $b_i + n$  as each of  $n$  partitions may be partially full
- Writing them back requires

# Cost of Hash Join

- Initial reading of a relation  $i$  requires  $b_i$  transfers
- Total number of blocks after partitioning is at most  $b_i + n$  as each of  $n$  partitions may be partially full
- Writing them back requires  $(b_i + n)$  transfers
- Reading them again during matching requires

# Cost of Hash Join

- Initial reading of a relation  $i$  requires  $b_i$  transfers
- Total number of blocks after partitioning is at most  $b_i + n$  as each of  $n$  partitions may be partially full
- Writing them back requires  $(b_i + n)$  transfers
- Reading them again during matching requires  $(b_i + n)$  transfers
- Therefore, total number of transfers is  $3(b_r + b_s) + 4n$

# Cost of Hash Join

- Initial reading of a relation  $i$  requires  $b_i$  transfers
- Total number of blocks after partitioning is at most  $b_i + n$  as each of  $n$  partitions may be partially full
- Writing them back requires  $(b_i + n)$  transfers
- Reading them again during matching requires  $(b_i + n)$  transfers
- Therefore, total number of transfers is  $3(b_r + b_s) + 4n$
- Assume memory buffer of  $m$  blocks
- Partitioning requires



# Cost of Hash Join

- Initial reading of a relation  $i$  requires  $b_i$  transfers
- Total number of blocks after partitioning is at most  $b_i + n$  as each of  $n$  partitions may be partially full
- Writing them back requires  $(b_i + n)$  transfers
- Reading them again during matching requires  $(b_i + n)$  transfers
- Therefore, total number of transfers is  $3(b_r + b_s) + 4n$
- Assume memory buffer of  $m$  blocks
- Partitioning requires  $k_i = \lceil b_i/m \rceil$  seeks for reading and  $k'_i = \lceil (b_i + n)/m \rceil$  seeks for writing
- Reading  $n$  partitions during matching requires  $n$  seeks per relation
- Therefore, total number of seeks is  $k_r + k_s + k'_r + k'_s + 2n$

# Cost of Recursive Partitioning

- Number of passes in recursive partitioning is

# Cost of Recursive Partitioning

- Number of passes in recursive partitioning is  $p = \lfloor \log_m b_s \rfloor$

# Cost of Recursive Partitioning

- Number of passes in recursive partitioning is  $p = \lfloor \log_m b_s \rfloor$
- During partitioning, total number of block transfers is

# Cost of Recursive Partitioning

- Number of passes in recursive partitioning is  $p = \lfloor \log_m b_s \rfloor$
- During partitioning, total number of block transfers is  $2p(b_r + n + b_s + n)$

# Cost of Recursive Partitioning

- Number of passes in recursive partitioning is  $p = \lfloor \log_m b_s \rfloor$
- During partitioning, total number of block transfers is  $2p(b_r + n + b_s + n)$
- Number of seeks during partitioning is

# Cost of Recursive Partitioning


- Number of passes in recursive partitioning is  $p = \lfloor \log_m b_s \rfloor$
- During partitioning, total number of block transfers is  $2p(b_r + n + b_s + n)$
- Number of seeks during partitioning is  $2p(k'_r + k'_s)$

# Cost of Recursive Partitioning

- Number of passes in recursive partitioning is  $p = \lfloor \log_m b_s \rfloor$
- During partitioning, total number of block transfers is  $2p(b_r + n + b_s + n)$
- Number of seeks during partitioning is  $2p(k'_r + k'_s)$
- Matching requires



# Cost of Recursive Partitioning

- Number of passes in recursive partitioning is  $p = \lfloor \log_m b_s \rfloor$
- During partitioning, total number of block transfers is  $2p(b_r + n + b_s + n)$
- Number of seeks during partitioning is  $2p(k'_r + k'_s)$  
- Matching requires  $(b_r + n + b_s + n)$  transfers and  $n + n$  seeks

# Example

- $r$ : 40000 tuples at 200 per block:

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block:

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 200/15 \rceil = 14$
- Size of each partition of  $r$  is



# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 200/15 \rceil = 14$
- Size of each partition of  $r$  is  $\lceil 250/15 \rceil = 17$

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 200/15 \rceil = 14$
- Size of each partition of  $r$  is  $\lceil 250/15 \rceil = 17$
- Build input should be

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 250/15 \rceil = 17$
- Size of each partition of  $r$  is  $\lceil 200/15 \rceil = 14$
- Build input should be  $s$

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 250/15 \rceil = 17$
- Size of each partition of  $r$  is  $\lceil 200/15 \rceil = 14$
- Build input should be  $s$
- Partitioning  $s$  requires

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 250/15 \rceil = 17$
- Size of each partition of  $r$  is  $\lceil 200/15 \rceil = 14$
- Build input should be  $s$
- Partitioning  $s$  requires  $200 + (200 + 15) = 415$  transfers
- Number of seeks for  $s$  is

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 200/15 \rceil = 14$
- Size of each partition of  $r$  is  $\lceil 250/15 \rceil = 17$
- Build input should be  $s$
- Partitioning  $s$  requires  $200 + (200 + 15) = 415$  transfers
- Number of seeks for  $s$  is  $\lceil 200/10 \rceil + \lceil 215/10 \rceil = 42$

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 200/15 \rceil = 14$
- Size of each partition of  $r$  is  $\lceil 250/15 \rceil = 17$
- Build input should be  $s$
- Partitioning  $s$  requires  $200 + (200 + 15) = 415$  transfers
- Number of seeks for  $s$  is  $\lceil 200/10 \rceil + \lceil 215/10 \rceil = 42$
- Partitioning  $r$  requires

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 200/15 \rceil = 14$
- Size of each partition of  $r$  is  $\lceil 250/15 \rceil = 17$
- Build input should be  $s$
- Partitioning  $s$  requires  $200 + (200 + 15) = 415$  transfers
- Number of seeks for  $s$  is  $\lceil 200/10 \rceil + \lceil 215/10 \rceil = 42$
- Partitioning  $r$  requires 1 pass since it is probe input
- Number of transfers for  $r$  is



# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 200/15 \rceil = 14$
- Size of each partition of  $r$  is  $\lceil 250/15 \rceil = 17$
- Build input should be  $s$
- Partitioning  $s$  requires  $200 + (200 + 15) = 415$  transfers
- Number of seeks for  $s$  is  $\lceil 200/10 \rceil + \lceil 215/10 \rceil = 42$
- Partitioning  $r$  requires 1 pass since it is probe input
- Number of transfers for  $r$  is  $250 + (250 + 15) = 515$
- Number of seeks for  $r$  is

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 200/15 \rceil = 14$
- Size of each partition of  $r$  is  $\lceil 250/15 \rceil = 17$
- Build input should be  $s$
- Partitioning  $s$  requires  $200 + (200 + 15) = 415$  transfers
- Number of seeks for  $s$  is  $\lceil 200/10 \rceil + \lceil 215/10 \rceil = 42$
- Partitioning  $r$  requires 1 pass since it is probe input
- Number of transfers for  $r$  is  $250 + (250 + 15) = 515$
- Number of seeks for  $r$  is  $\lceil 250/10 \rceil + \lceil 265/10 \rceil = 52$

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 200/15 \rceil = 14$
- Size of each partition of  $r$  is  $\lceil 250/15 \rceil = 17$
- Build input should be  $s$
- Partitioning  $s$  requires  $200 + (200 + 15) = 415$  transfers
- Number of seeks for  $s$  is  $\lceil 200/10 \rceil + \lceil 215/10 \rceil = 42$
- Partitioning  $r$  requires 1 pass since it is probe input
- Number of transfers for  $r$  is  $250 + (250 + 15) = 515$
- Number of seeks for  $r$  is  $\lceil 250/10 \rceil + \lceil 265/10 \rceil = 52$
- Matching phase requires reading all blocks of  $s$  and  $r$ :

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 200/15 \rceil = 14$
- Size of each partition of  $r$  is  $\lceil 250/15 \rceil = 17$
- Build input should be  $s$
- Partitioning  $s$  requires  $200 + (200 + 15) = 415$  transfers
- Number of seeks for  $s$  is  $\lceil 200/10 \rceil + \lceil 215/10 \rceil = 42$
- Partitioning  $r$  requires 1 pass since it is probe input
- Number of transfers for  $r$  is  $250 + (250 + 15) = 515$
- Number of seeks for  $r$  is  $\lceil 250/10 \rceil + \lceil 265/10 \rceil = 52$
- Matching phase requires reading all blocks of  $s$  and  $r$ :  
 $215 + 265 = 480$
- Number of seeks in matching phase is

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 200/15 \rceil = 14$
- Size of each partition of  $r$  is  $\lceil 250/15 \rceil = 17$
- Build input should be  $s$
- Partitioning  $s$  requires  $200 + (200 + 15) = 415$  transfers
- Number of seeks for  $s$  is  $\lceil 200/10 \rceil + \lceil 215/10 \rceil = 42$
- Partitioning  $r$  requires 1 pass since it is probe input
- Number of transfers for  $r$  is  $250 + (250 + 15) = 515$
- Number of seeks for  $r$  is  $\lceil 250/10 \rceil + \lceil 265/10 \rceil = 52$
- Matching phase requires reading all blocks of  $s$  and  $r$ :  
 $215 + 265 = 480$
- Number of seeks in matching phase is  $15 + 15 = 30$

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 200/15 \rceil = 14$
- Size of each partition of  $r$  is  $\lceil 250/15 \rceil = 17$
- Build input should be  $s$
- Partitioning  $s$  requires  $200 + (200 + 15) = 415$  transfers
- Number of seeks for  $s$  is  $\lceil 200/10 \rceil + \lceil 215/10 \rceil = 42$
- Partitioning  $r$  requires 1 pass since it is probe input
- Number of transfers for  $r$  is  $250 + (250 + 15) = 515$
- Number of seeks for  $r$  is  $\lceil 250/10 \rceil + \lceil 265/10 \rceil = 52$
- Matching phase requires reading all blocks of  $s$  and  $r$ :  
 $215 + 265 = 480$
- Number of seeks in matching phase is  $15 + 15 = 30$
- Total is

# Example

- $r$ : 40000 tuples at 200 per block: 200 blocks
- $s$ : 25000 tuples at 100 per block: 250 blocks
- Available memory is 15 for partitions and  $m = 10$  for buffering
- Number of partitions is  $n = 15$
- Size of each partition of  $s$  is  $\lceil 200/15 \rceil = 14$
- Size of each partition of  $r$  is  $\lceil 250/15 \rceil = 17$
- Build input should be  $s$
- Partitioning  $s$  requires  $200 + (200 + 15) = 415$  transfers
- Number of seeks for  $s$  is  $\lceil 200/10 \rceil + \lceil 215/10 \rceil = 42$
- Partitioning  $r$  requires 1 pass since it is probe input
- Number of transfers for  $r$  is  $250 + (250 + 15) = 515$
- Number of seeks for  $r$  is  $\lceil 250/10 \rceil + \lceil 265/10 \rceil = 52$
- Matching phase requires reading all blocks of  $s$  and  $r$ :  
 $215 + 265 = 480$
- Number of seeks in matching phase is  $15 + 15 = 30$
- Total is  $15 + 15 + 42 + 52 = 124$  seeks and  
 $415 + 515 + 480 = 1410$  transfers

# Other Operations

- Complex join conditions



# Other Operations

- Complex join conditions
  - Block nested-loop for conjunctive and/or disjunctive selection
  - If only equality, union/intersection of results of index or merge or hash join may be used

# Other Operations

- Complex join conditions
  - Block nested-loop for conjunctive and/or disjunctive selection
  - If only equality, union/intersection of results of index or merge or hash join may be used
- Outer join

# Other Operations

- Complex join conditions
  - Block nested-loop for conjunctive and/or disjunctive selection
  - If only equality, union/intersection of results of index or merge or hash join may be used
- Outer join
  - Block nested-loop algorithm requires almost no modification

# Other Operations

- Complex join conditions
  - Block nested-loop for conjunctive and/or disjunctive selection
  - If only equality, union/intersection of results of index or merge or hash join may be used
- Outer join
  - Block nested-loop algorithm requires almost no modification
  - For merge join, select all (non-joining) records while scanning

# Other Operations

- Complex join conditions
  - Block nested-loop for conjunctive and/or disjunctive selection
  - If only equality, union/intersection of results of index or merge or hash join may be used
- Outer join
  - Block nested-loop algorithm requires almost no modification
  - For merge join, select all (non-joining) records while scanning
  - For hash join, if  $r \bowtie s$ ,  $r$  should be the

# Other Operations

- Complex join conditions
  - Block nested-loop for conjunctive and/or disjunctive selection
  - If only equality, union/intersection of results of index or merge or hash join may be used
- Outer join
  - Block nested-loop algorithm requires almost no modification
  - For merge join, select all (non-joining) records while scanning
  - For hash join, if  $r \supsetneq s$ ,  $r$  should be the probe relation
  - If  $r$  is the build relation, keep track of which records in hash index have been used; output all non-used records
  - If  $r \supsetneq s$ , use both techniques

# Other Operations

- Complex join conditions
  - Block nested-loop for conjunctive and/or disjunctive selection
  - If only equality, union/intersection of results of index or merge or hash join may be used
- Outer join
  - Block nested-loop algorithm requires almost no modification
  - For merge join, select all (non-joining) records while scanning
  - For hash join, if  $r \supsetneq s$ ,  $r$  should be the probe relation
  - If  $r$  is the build relation, keep track of which records in hash index have been used; output all non-used records
  - If  $r \supsetneq s$ , use both techniques
- Set operations

# Other Operations

- Complex join conditions
  - Block nested-loop for conjunctive and/or disjunctive selection
  - If only equality, union/intersection of results of index or merge or hash join may be used
- Outer join
  - Block nested-loop algorithm requires almost no modification
  - For merge join, select all (non-joining) records while scanning
  - For hash join, if  $r \bowtie s$ ,  $r$  should be the probe relation
  - If  $r$  is the build relation, keep track of which records in hash index have been used; output all non-used records
  - If  $r \bowtie s$ , use both techniques
- Set operations
  - If relations are sorted, scan in order
  - Build hash index on one relation; test records from other relation



# Other Operations

- Complex join conditions
  - Block nested-loop for conjunctive and/or disjunctive selection
  - If only equality, union/intersection of results of index or merge or hash join may be used
- Outer join
  - Block nested-loop algorithm requires almost no modification
  - For merge join, select all (non-joining) records while scanning
  - For hash join, if  $r \supsetneq s$ ,  $r$  should be the probe relation
  - If  $r$  is the build relation, keep track of which records in hash index have been used; output all non-used records
  - If  $r \supsetneq s$ , use both techniques
- Set operations
  - If relations are sorted, scan in order
  - Build hash index on one relation; test records from other relation
- Grouping and aggregation

# Other Operations

- Complex join conditions
  - Block nested-loop for conjunctive and/or disjunctive selection
  - If only equality, union/intersection of results of index or merge or hash join may be used
- Outer join
  - Block nested-loop algorithm requires almost no modification
  - For merge join, select all (non-joining) records while scanning
  - For hash join, if  $r \supsetneq s$ ,  $r$  should be the probe relation
  - If  $r$  is the build relation, keep track of which records in hash index have been used; output all non-used records
  - If  $r \supsetneq s$ , use both techniques
- Set operations
  - If relations are sorted, scan in order
  - Build hash index on one relation; test records from other relation
- Grouping and aggregation
  - Use hashing to organize into groups; then apply aggregation

# Other Operations

- Complex join conditions
  - Block nested-loop for conjunctive and/or disjunctive selection
  - If only equality, union/intersection of results of index or merge or hash join may be used
- Outer join
  - Block nested-loop algorithm requires almost no modification
  - For merge join, select all (non-joining) records while scanning
  - For hash join, if  $r \supsetneq s$ ,  $r$  should be the probe relation
  - If  $r$  is the build relation, keep track of which records in hash index have been used; output all non-used records
  - If  $r \supsetneq s$ , use both techniques
- Set operations
  - If relations are sorted, scan in order
  - Build hash index on one relation; test records from other relation
- Grouping and aggregation
  - Use hashing to organize into groups; then apply aggregation
- Duplicate detection and elimination

# Other Operations

- Complex join conditions
  - Block nested-loop for conjunctive and/or disjunctive selection
  - If only equality, union/intersection of results of index or merge or hash join may be used
- Outer join
  - Block nested-loop algorithm requires almost no modification
  - For merge join, select all (non-joining) records while scanning
  - For hash join, if  $r \supsetneq s$ ,  $r$  should be the probe relation
  - If  $r$  is the build relation, keep track of which records in hash index have been used; output all non-used records
  - If  $r \supsetneq s$ , use both techniques
- Set operations
  - If relations are sorted, scan in order
  - Build hash index on one relation; test records from other relation
- Grouping and aggregation
  - Use hashing to organize into groups; then apply aggregation
- Duplicate detection and elimination
  - Use hashing or sorting