

# CS315: DATABASE SYSTEMS NoSQL AND BIG DATA SYSTEMS

Arnab Bhattacharya

`arnabb@cse.iitk.ac.in`

Computer Science and Engineering,  
Indian Institute of Technology, Kanpur

`http://web.cse.iitk.ac.in/~cs315/`

2<sup>nd</sup> semester, 2022-23

Tue 10:30-11:45, Thu 12:00-13:15

- **NoSQL** aims to provide
  - Scalability
  - Flexibility
  - Naturalness
  - Distribution
  - Performance

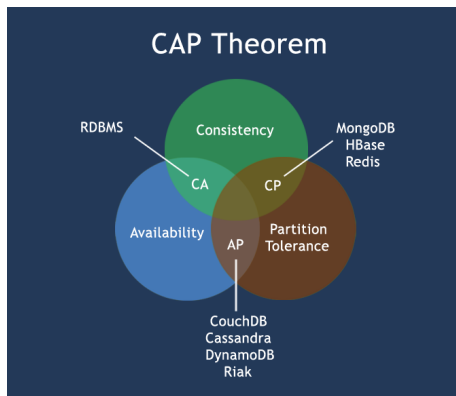
- **NoSQL** aims to provide
  - Scalability
  - Flexibility
  - Naturalness
  - Distribution
  - Performance
- NoSQL is not so much worried about the ACID properties
- Especially, consistency and atomicity are relaxed
- So, NoSQL started as “no-SQL”

- **NoSQL** aims to provide
  - Scalability
  - Flexibility
  - Naturalness
  - Distribution
  - Performance
- NoSQL is not so much worried about the ACID properties
- Especially, consistency and atomicity are relaxed
- So, NoSQL started as “no-SQL”
- Later changed since RDBMS is too powerful to always ignore

- **NoSQL** aims to provide
  - Scalability
  - Flexibility
  - Naturalness
  - Distribution
  - Performance
- NoSQL is not so much worried about the ACID properties
- Especially, consistency and atomicity are relaxed
- So, NoSQL started as “no-SQL”
- Later changed since RDBMS is too powerful to always ignore
- Now, NoSQL is **not only SQL**

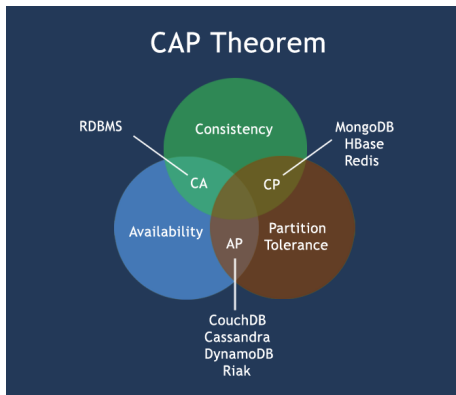
- **NoSQL** aims to provide
  - Scalability
  - Flexibility
  - Naturalness
  - Distribution
  - Performance
- NoSQL is not so much worried about the ACID properties
- Especially, consistency and atomicity are relaxed
- So, NoSQL started as “no-SQL”
- Later changed since RDBMS is too powerful to always ignore
- Now, NoSQL is **not only SQL**
- **NewSQL**

# CAP Theorem



- CA: single-site; partitioning is not allowed
- CP: what is available is consistent
- AP: everything is available but may not be consistent

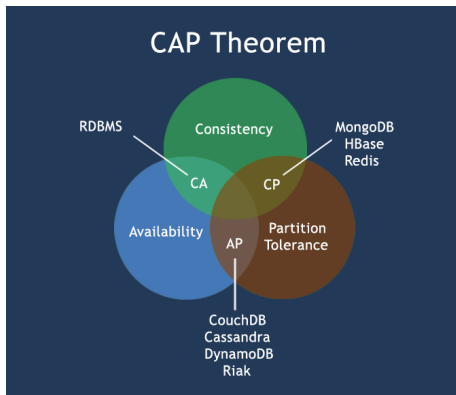
# CAP Theorem



- CA: single-site; partitioning is not allowed
- CP: what is available is consistent
- AP: everything is available but may not be consistent
- All of C, A, P **cannot** be satisfied simultaneously



# CAP Theorem



- CA: single-site; partitioning is not allowed
- CP: what is available is consistent
- AP: everything is available but may not be consistent
- All of C, A, P **cannot** be satisfied simultaneously
- More a hypothesis than a theorem

# BASE Properties

- **BASE properties** are to counter ACID

# BASE Properties

- **BASE properties** are to counter ACID
- **Basically Available**: System guarantees availability

# BASE Properties

- **BASE properties** are to counter ACID
- **Basically Available**: System guarantees availability
- **Soft state**: State of system is soft, i.e., it may change without input to maintain consistency

# BASE Properties

- **BASE properties** are to counter ACID
- **Basically Available**: System guarantees availability
- **Soft state**: State of system is soft, i.e., it may change without input to maintain consistency
- **Eventual consistency**: Data will be eventually consistent without any interim perturbation
  - Sacrifices strong (immediate) consistency

# Types

- Main types of NoSQL data stores:

- 1 Columnar families
- 2 Key-value stores
- 3 Bigtable systems
- 4 Document databases
- 5 Graph databases

# Columnar Storage

- Instead of rows being stored together, columns are stored consecutively
- A single disk block (or a set of consecutive blocks) stores a single **column family**
- A column family may consist of one or multiple columns
- This set of columns is called a **super column**

# Columnar Storage

- Instead of rows being stored together, columns are stored consecutively
- A single disk block (or a set of consecutive blocks) stores a single **column family**
- A column family may consist of one or multiple columns
- This set of columns is called a **super column**
- Two main types
  - Columnar relational models
  - Key-value stores and/or big tables



# Columnar Relational Models

- A type of *RDBMS*
- Column-wise storage on the disk

# Columnar Relational Models

- A type of *RDBMS*
- Column-wise storage on the disk
- Allows faster querying when only few columns are touched on the entire data
- Allows compression of columns
- Provides better memory caching
- Joins are faster since they are mostly on similar columns from two tables

# Columnar Relational Models

- A type of *RDBMS*
- Column-wise storage on the disk
- Allows faster querying when only few columns are touched on the entire data
- Allows compression of columns
- Provides better memory caching
- Joins are faster since they are mostly on similar columns from two tables
- Is not good for updates
- Is not good when many columns of a few tuples are accessed

# Columnar Relational Models

- A type of *RDBMS*
- Column-wise storage on the disk
- Allows faster querying when only few columns are touched on the entire data
- Allows compression of columns
- Provides better memory caching
- Joins are faster since they are mostly on similar columns from two tables
- Is not good for updates
- Is not good when many columns of a few tuples are accessed
- Good for **OLAP** (online analytical processing)
- Not good for **OLTP** (online transaction processing)

# Columnar Relational Models

- A type of *RDBMS*
- Column-wise storage on the disk
- Allows faster querying when only few columns are touched on the entire data
- Allows compression of columns
- Provides better memory caching
- Joins are faster since they are mostly on similar columns from two tables
- Is not good for updates
- Is not good when many columns of a few tuples are accessed
- Good for **OLAP** (online analytical processing)
- Not good for **OLTP** (online transaction processing)
- Example: MonetDB

# Key-Value Stores

- Two columns: a **key** and a **value**
- Key is mostly text
- Value can be anything and is simply an object

# Key-Value Stores

- Two columns: a **key** and a **value**
- Key is mostly text
- Value can be anything and is simply an object
- Essentially, actual data becomes “value” and an unique id is generated which becomes “key”

# Key-Value Stores

- Two columns: a **key** and a **value**
- Key is mostly text
- Value can be anything and is simply an object
- Essentially, actual data becomes “value” and an unique id is generated which becomes “key”
- Whole database is then just one big table with these two columns
- Becomes schema-less



# Key-Value Stores

- Two columns: a **key** and a **value**
- Key is mostly text
- Value can be anything and is simply an object
- Essentially, actual data becomes “value” and an unique id is generated which becomes “key”
- Whole database is then just one big table with these two columns
- Becomes schema-less
- Can be distributed and is, thus, highly scalable
- So, in essence a big distributed hash table

# Key-Value Stores

- Two columns: a **key** and a **value**
- Key is mostly text
- Value can be anything and is simply an object
- Essentially, actual data becomes “value” and an unique id is generated which becomes “key”
- Whole database is then just one big table with these two columns
- Becomes schema-less
- Can be distributed and is, thus, highly scalable
- So, in essence a big distributed hash table
- All queries are on keys
- Keys are necessarily indexed

# Key-Value Stores

- Two columns: a **key** and a **value**
- Key is mostly text
- Value can be anything and is simply an object
- Essentially, actual data becomes “value” and an unique id is generated which becomes “key”
- Whole database is then just one big table with these two columns
- Becomes schema-less
- Can be distributed and is, thus, highly scalable
- So, in essence a big distributed hash table
- All queries are on keys
- Keys are necessarily indexed
- Example: Cassandra, CouchDB, Tokyo Cabinet, Redis

# BigTable Systems

- Started from Google's BigTable implementation
- Uses a key-value store
- Data can be replicated for better availability

# BigTable Systems

- Started from Google's BigTable implementation
- Uses a key-value store
- Data can be replicated for better availability
- Uses a **timestamp**
- Timestamp is used to
  - Expire data
  - Delete stale data
  - Resolve read-write conflicts

# BigTable Systems

- Started from Google's BigTable implementation
- Uses a key-value store
- Data can be replicated for better availability
- Uses a **timestamp**
- Timestamp is used to
  - Expire data
  - Delete stale data
  - Resolve read-write conflicts
- Same value can be indexed using multiple keys
- Map-reduce framework to compute

# BigTable Systems

- Started from Google's BigTable implementation
- Uses a key-value store
- Data can be replicated for better availability
- Uses a **timestamp**
- Timestamp is used to
  - Expire data
  - Delete stale data
  - Resolve read-write conflicts
- Same value can be indexed using multiple keys
- Map-reduce framework to compute
- Example: BigTable, HBase, Cassandra, HyperTable, SimpleDB

# Document Databases

- Uses document structure as the main storage format of data
- Popular document formats are XML, JSON, BSON, YAML
- Document itself is the key while the content is the value
- Document can be indexed by id or simply its location (e.g., URI)



# Document Databases

- Uses document structure as the main storage format of data
- Popular document formats are XML, JSON, BSON, YAML
- Document itself is the key while the content is the value
- Document can be indexed by id or simply its location (e.g., URI)
- Content needs to be parsed to make sense
- Content can be organised further

# Document Databases

- Uses document structure as the main storage format of data
- Popular document formats are XML, JSON, BSON, YAML
- Document itself is the key while the content is the value
- Document can be indexed by id or simply its location (e.g., URI)
- Content needs to be parsed to make sense
- Content can be organised further
- Extremely useful for insert-once read-many scenarios
- Can use map-reduce framework to compute

# Document Databases

- Uses document structure as the main storage format of data
- Popular document formats are XML, JSON, BSON, YAML
- Document itself is the key while the content is the value
- Document can be indexed by id or simply its location (e.g., URI)
- Content needs to be parsed to make sense
- Content can be organised further
- Extremely useful for insert-once read-many scenarios
- Can use map-reduce framework to compute
- Example: MongoDB, CouchDB

# Graph Databases

- Nodes represent entities
- Edges encode relationships between nodes
- Can be directed
- Can have hyper-edges as well

# Graph Databases

- Nodes represent entities
- Edges encode relationships between nodes
- Can be directed
- Can have hyper-edges as well
- Easier to find distances and neighbors

# Graph Databases

- Nodes represent entities
- Edges encode relationships between nodes
- Can be directed
- Can have hyper-edges as well
- Easier to find distances and neighbors
- Example: Neo4J, HyperGraph, Infinite Graph, Titan, FlockDB

# Discussion

- NoSQL, although started as anti-SQL, is no more so
- More a realisation that for some cases
  - RDBMS does not scale or distribute, or
  - ACIDity is an overkill

# Discussion

- NoSQL, although started as anti-SQL, is no more so
- More a realisation that for some cases
  - RDBMS does not scale or distribute, or
  - ACIDity is an overkill
- NoSQL is *not* good for every scenario
  - Not always that strong consistency can be sacrificed
- Most legacy systems still use RDBMS



# Discussion

- NoSQL, although started as anti-SQL, is no more so
- More a realisation that for some cases
  - RDBMS does not scale or distribute, or
  - ACIDity is an overkill
- NoSQL is *not* good for every scenario
  - Not always that strong consistency can be sacrificed
- Most legacy systems still use RDBMS
- NoSQL horizon is shifting rapidly
- Trend is for NoSQL as cloud computing and big data relies on it

# Discussion

- NoSQL, although started as anti-SQL, is no more so
- More a realisation that for some cases
  - RDBMS does not scale or distribute, or
  - ACIDity is an overkill
- NoSQL is *not* good for every scenario
  - Not always that strong consistency can be sacrificed
- Most legacy systems still use RDBMS
- NoSQL horizon is shifting rapidly
- Trend is for NoSQL as cloud computing and big data relies on it
- Many NoSQL systems are increasingly using features of RDBMS
- New paradigm of scalability with transaction support is **NewSQL**

# Big Data

- What is **big data**?

# Big Data

- What is **big data**?
- Data which is big
- How big is “big”?

# Big Data

- What is **big data**?
- Data which is big
- How big is “big”?
- For sociologists, 100 subjects is big

# Big Data

- What is **big data**?
- Data which is big
- How big is “big”?
- For sociologists, 100 subjects is big
- For social networks, terrabytes is normal

# Big Data

- What is **big data**?
- Data which is big
- How big is “big”?
- For sociologists, 100 subjects is big
- For social networks, terrabytes is normal
- For astrophysicists, terrabytes is an hour's work

# Big Data

- What is **big data**?
- Data which is big
- How big is “big”?
- For sociologists, 100 subjects is big
- For social networks, terrabytes is normal
- For astrophysicists, terrabytes is an hour’s work
- So, no absolute definition or threshold



# Big Data

- What is **big data**?
- Data which is big
- How big is “big”?
- For sociologists, 100 subjects is big
- For social networks, terrabytes is normal
- For astrophysicists, terrabytes is an hour’s work
- So, no absolute definition or threshold
- When data is bigger than most standard machines can store or most algorithms can handle


# Characterization of Big Data

- Having large volumes of data requires
  - Newer techniques
  - Newer tools
  - Newer architectures

# Characterization of Big Data


- Having large volumes of data requires
  - Newer techniques
  - Newer tools
  - Newer architectures
- Allows solving newer problems
  - Can also solve older problems better

# Properties of Big Data



- 3 V's: **volume**, **variety**, **velocity**
- **Volume**: When data is extremely large in size, how to load it, index it or query it
- **Variety**: Data can be semi-structured or unstructured as well; how to query
- **Velocity**: Data can arrive at real time and can be streaming

# Properties of Big Data



- 3 V's: **volume, variety, velocity**
- **Volume**: When data is extremely large in size, how to load it, index it or query it
- **Variety**: Data can be semi-structured or unstructured as well; how to query
- **Velocity**: Data can arrive at real time and can be streaming
- Extended V's: **veracity, validity, visibility, variability**

# Enablers of Big Data

- Improved storage volume and type
- Improved processing power
- Improved data
- Improved network speed
- Improved architecture

# Enablers of Big Data

- Improved storage volume and type
- Improved processing power
- Improved data
- Improved network speed
- Improved architecture
- Increased capital
- Increased business

# Tools for Big Data

- Hosting: Distributed servers or Cloud
  - Amazon EC2



# Tools for Big Data

- Hosting: Distributed servers or Cloud
  - Amazon EC2
- File system: Scalable and distributed
  - HDFS, Amazon S3

# Tools for Big Data

- Hosting: Distributed servers or Cloud
  - Amazon EC2
- File system: Scalable and distributed
  - HDFS, Amazon S3
- Programming model: Distributed scalable processing
  - Map-reduce framework, Hadoop

# Tools for Big Data

- Hosting: Distributed servers or Cloud
  - Amazon EC2
- File system: Scalable and distributed
  - HDFS, Amazon S3
- Programming model: Distributed scalable processing
  - Map-reduce framework, Hadoop
- Database: NoSQL
  - HBase, MongoDB, Cassandra

# Tools for Big Data

- Hosting: Distributed servers or Cloud
  - Amazon EC2
- File system: Scalable and distributed
  - HDFS, Amazon S3
- Programming model: Distributed scalable processing
  - Map-reduce framework, Hadoop
- Database: NoSQL
  - HBase, MongoDB, Cassandra
- Operations: Querying, indexing, analytics
  - Data mining, Information retrieval
  - Machine learning: Mahout on top of Hadoop

# Discussion

- Not so new term: **Data Science (DS)**
  - Connection with **Machine Learning (ML)** and **Artificial Intelligence (AI)**

# Discussion

- Not so new term: **Data Science (DS)**
  - Connection with **Machine Learning (ML)** and **Artificial Intelligence (AI)**
- Many open-source tools
- Cloud, etc. can be rented

# Discussion

- Not so new term: **Data Science (DS)**
  - Connection with **Machine Learning (ML)** and **Artificial Intelligence (AI)**
- Many open-source tools
- Cloud, etc. can be rented
- Many applications still do *not* require big data

# Discussion

- Not so new term: **Data Science (DS)**
  - Connection with **Machine Learning (ML)** and **Artificial Intelligence (AI)**
- Many open-source tools
- Cloud, etc. can be rented
- Many applications still do *not* require big data
- Databases have become robust and stable